

## ELF – Class Notes

- Executable binary files are constructed or build using a binary file format standard.  
E.g. ELF
- Most of the binary file formats organize executable image in the binary file as per the platforms memory model.
- Executable file header provides detail about the platform for which the file is built and information about sections and segments that are part of executable image.
- To view the ELF header file of any executable, use a tool called **readelf**. Let's consider one executable image **app**.

**\$readelf -a app | more**

```

veda@linux: ~/dl
ELF Header:
  Magic:   7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
  Class:                   ELF32
  Data:                   2's complement, little endian
  Version:               1 (current)
  OS/ABI:                UNIX - System V
  ABI Version:           0
  Type:                  EXEC (Executable file)
  Machine:               Intel 80386
  Version:               0x1
  Entry point address:    0x8048470
  Start of program headers: 52 (bytes into file)
  Start of section headers: 4416 (bytes into file)
  Flags:                  0x0
  Size of this header:    52 (bytes)
  Size of program headers: 32 (bytes)
  Number of program headers: 8
  Size of section headers: 40 (bytes)
  Number of section headers: 29
  Section header string table index: 26

Section Headers:
[Nr] Name                Type              Addr      Off      Size    ES Flg  Lk  Inf Al
[ 0]                     NULL              00000000  000000  000000  00   0  0  0  0
[ 1] .interp               PROGBITS          08048134  000134  000013  00   A  0  0  1
[ 2] .note.ABI-tag         NOTE              08048148  000148  000020  00   A  0  0  4
[ 3] .note.gnu.build-id    NOTE              08048168  000168  000024  00   A  0  0  4
[ 4] .gnu.hash             GNU_HASH          0804818c  00018c  000020  04   A  5  0  4
[ 5] .dynsym               DYNSYM            080481ac  0001ac  0000b0  10   A  6  1  4
[ 6] .dynstr               STRTAB            0804825c  00025c  000097  00   A  0  0  1
[ 7] .gnu.version          VERSYM            080482f4  0002f4  000016  02   A  5  0  2
[ 8] .gnu.version_r        VERNEED           0804830c  00030c  000050  00   A  6  2  4
[ 9] .rel.dyn              REL               0804835c  00035c  000008  08   A  5  0  4
[10] .rel.plt              REL               08048364  000364  000040  08   A  5 12  4
[11] .init                 PROGBITS          080483a4  0003a4  000030  00  AX  0  0  4
[12] .plt                 PROGBITS          080483d4  0003d4  000090  04  AX  0  0  4
[13] .text                PROGBITS          08048470  000470  0001ec  00  AX  0  0 16
[14] .fini                PROGBITS          0804865c  00065c  00001c  00  AX  0  0  4
[15] .rodata              PROGBITS          08048678  000678  000033  00   A  0  0  4
[16] .eh_frame            PROGBITS          080486ac  0006ac  000004  00   A  0  0  4
[17] .ctors               PROGBITS          08049f0c  000f0c  000008  00  WA  0  0  4
[18] .dtors               PROGBITS          08049f14  000f14  000008  00  WA  0  0  4
--More--
  
```

Elf header for an executable provides following important details

- Platform's application binary interface standard considered while building the executable. ABI standard specifies the following:
  - Addressing format to be used by assembler and linker while binding machine instructions to a corresponding address
  - Function calling conventions (fast call , standard call , windows calling conventions)
  - Binary file format to be used to build object files
- It specifies object file type, mainly three types relocatable, executable, shared object.
- It specifies base address or entry point address of code segment (usually base address is `_start`)
- Offset of the program header table in the file

### Load time and Run time Libraries

#### Load time Libraries:

- When a shared object is linked into application address space during application initialization, it is referred as a load time library.
- All symbols linked into executables from a dynamic library are resolved during application initialization.
- Load time libraries remain resident in the process until application terminates.

#### Run time Libraries:

- Applications at run time could raise a request to load shared objects and bind them into address space. Such shared objects can also be detached again if application makes a request.
- If an application intends to use a shared object as runtime library, all direct references to library symbol must be avoided.

In order to understand **run time libraries** in detail, let's take a simple **C** program **test.c** which calls a shared library **libxyz.so**, where the function '**add**' resides.

**\$vim test.c**

```
veda@linux: ~/dl
#include<stdio.h>
#include<dlfcn.h>

main()
{
    char *ptr;
    int i;
    void *handle;
    int (*fnptr)(int , int);
    getchar();
    handle=dlopen("./libxyz.so",RTLD_NOW);
    if(handle==NULL) {
        printf("\n Failed \n");
    }
    fnptr = dlsym(handle,"add"); /* it returns the address of the function */
    getchar();
    i=(fnptr)(20,20);
    printf("\n the result: %d",i);
    dlclose(handle);
    getchar();
}
```

In the program **getchar()** is included to put the process under wait stage. Compile it by including **-ldl**, where the **dlopen** ,**dlsym** , **dlclose** are defined.

**\$gcc test.c -o test -ldl**

Execute the output, on the first **getchar()** . Get the **PID** of our process using **ps -Af** and check **\$cat /proc/2104/maps**.

It shows as below

```
veda@linux: ~/dl
veda@linux:~/dl$ cat /proc/2104/maps
00110000-0026c000 r-xp 00000000 08:01 18623683 /lib/i386-linux-gnu/libc-2.13.so
0026c000-0026e000 r--p 0015c000 08:01 18623683 /lib/i386-linux-gnu/libc-2.13.so
0026e000-0026f000 rw-p 0015e000 08:01 18623683 /lib/i386-linux-gnu/libc-2.13.so
0026f000-00272000 rw-p 00000000 00:00 0
004ed000-004ee000 r-xp 00000000 00:00 0
006fc000-00718000 r-xp 00000000 08:01 18623680 /lib/i386-linux-gnu/ld-2.13.so
00718000-00719000 r--p 0001b000 08:01 18623680 /lib/i386-linux-gnu/ld-2.13.so
00719000-0071a000 rw-p 0001c000 08:01 18623680 /lib/i386-linux-gnu/ld-2.13.so
00b64000-00b66000 r-xp 00000000 08:01 18623686 /lib/i386-linux-gnu/libdl-2.13.so
00b66000-00b67000 r--p 00001000 08:01 18623686 /lib/i386-linux-gnu/libdl-2.13.so
00b67000-00b68000 rw-p 00002000 08:01 18623686 /lib/i386-linux-gnu/libdl-2.13.so
08048000-08049000 r-xp 00000000 08:01 9045280 /home/veda/dl/dlopen
08049000-0804a000 r--p 00000000 08:01 9045280 /home/veda/dl/dlopen
0804a000-0804b000 rw-p 00001000 08:01 9045280 /home/veda/dl/dlopen
b7740000-b7742000 rw-p 00000000 00:00 0
b7754000-b7757000 rw-p 00000000 00:00 0
bfed5000-bfef6000 rw-p 00000000 00:00 0 [stack]
veda@linux:~/dl$
```

It shows our library **libxyz.so** has not loaded.

Execute the second **getchar()** and check the maps again.



```
veda@linux: ~/dl
veda@linux:~/dl$ cat /proc/2104/maps
00110000-0026c000 r-xp 00000000 08:01 18623683 /lib/i386-linux-gnu/libc-2.13.so
0026c000-0026e000 r--p 0015c000 08:01 18623683 /lib/i386-linux-gnu/libc-2.13.so
0026e000-0026f000 rw-p 0015e000 08:01 18623683 /lib/i386-linux-gnu/libc-2.13.so
0026f000-00272000 rw-p 00000000 00:00 0
004ed000-004ee000 r-xp 00000000 00:00 0 [vdso]
006fc000-00718000 r-xp 00000000 08:01 18623680 /lib/i386-linux-gnu/ld-2.13.so
00718000-00719000 r--p 0001b000 08:01 18623680 /lib/i386-linux-gnu/ld-2.13.so
00719000-0071a000 rw-p 0001c000 08:01 18623680 /lib/i386-linux-gnu/ld-2.13.so
00737000-00738000 r-xp 00000000 08:01 9045282 /home/veda/dl/libxyz.so
00738000-00739000 r--p 00000000 08:01 9045282 /home/veda/dl/libxyz.so
00739000-0073a000 rw-p 00001000 08:01 9045282 /home/veda/dl/libxyz.so
00b64000-00b66000 r-xp 00000000 08:01 18623686 /lib/i386-linux-gnu/libdl-2.13.so
00b66000-00b67000 r--p 00001000 08:01 18623686 /lib/i386-linux-gnu/libdl-2.13.so
00b67000-00b68000 rw-p 00002000 08:01 18623686 /lib/i386-linux-gnu/libdl-2.13.so
08048000-08049000 r-xp 00000000 08:01 9045280 /home/veda/dl/dlopen
08049000-0804a000 r--p 00000000 08:01 9045280 /home/veda/dl/dlopen
0804a000-0804b000 rw-p 00001000 08:01 9045280 /home/veda/dl/dlopen
088b8000-088d9000 rw-p 00000000 00:00 0 [heap]
b7740000-b7742000 rw-p 00000000 00:00 0
b7754000-b7757000 rw-p 00000000 00:00 0
bfed5000-bfef6000 rw-p 00000000 00:00 0 [stack]
veda@linux:~/dl$
```

Here we will find our library **libxyz.so** , which has loaded into the process.

Then execute the third **getchar()** and check maps again

```
veda@linux: ~/dl
veda@linux:~/dl$ cat /proc/2104/maps
00110000-0026c000 r-xp 00000000 08:01 18623683 /lib/i386-linux-gnu/libc-2.13.so
0026c000-0026e000 r--p 0015c000 08:01 18623683 /lib/i386-linux-gnu/libc-2.13.so
0026e000-0026f000 rw-p 0015e000 08:01 18623683 /lib/i386-linux-gnu/libc-2.13.so
0026f000-00272000 rw-p 00000000 00:00 0
004ed000-004ee000 r-xp 00000000 00:00 0 [vdso]
006fc000-00718000 r-xp 00000000 08:01 18623680 /lib/i386-linux-gnu/ld-2.13.so
00718000-00719000 r--p 0001b000 08:01 18623680 /lib/i386-linux-gnu/ld-2.13.so
00719000-0071a000 rw-p 0001c000 08:01 18623680 /lib/i386-linux-gnu/ld-2.13.so
00b64000-00b66000 r-xp 00000000 08:01 18623686 /lib/i386-linux-gnu/libdl-2.13.so
00b66000-00b67000 r--p 00001000 08:01 18623686 /lib/i386-linux-gnu/libdl-2.13.so
00b67000-00b68000 rw-p 00002000 08:01 18623686 /lib/i386-linux-gnu/libdl-2.13.so
08048000-08049000 r-xp 00000000 08:01 9045280 /home/veda/dl/dlopen
08049000-0804a000 r--p 00000000 08:01 9045280 /home/veda/dl/dlopen
0804a000-0804b000 rw-p 00001000 08:01 9045280 /home/veda/dl/dlopen
088b8000-088d9000 rw-p 00000000 00:00 0 [heap]
b7740000-b7742000 rw-p 00000000 00:00 0
b7753000-b7757000 rw-p 00000000 00:00 0
bfed5000-bfef6000 rw-p 00000000 00:00 0 [stack]
veda@linux:~/dl$
```

Here it doesn't shows our library which clearly explains the functionality of a run time library. It gets loaded when the program calls it and gets unloaded as soon as the function gets over.