


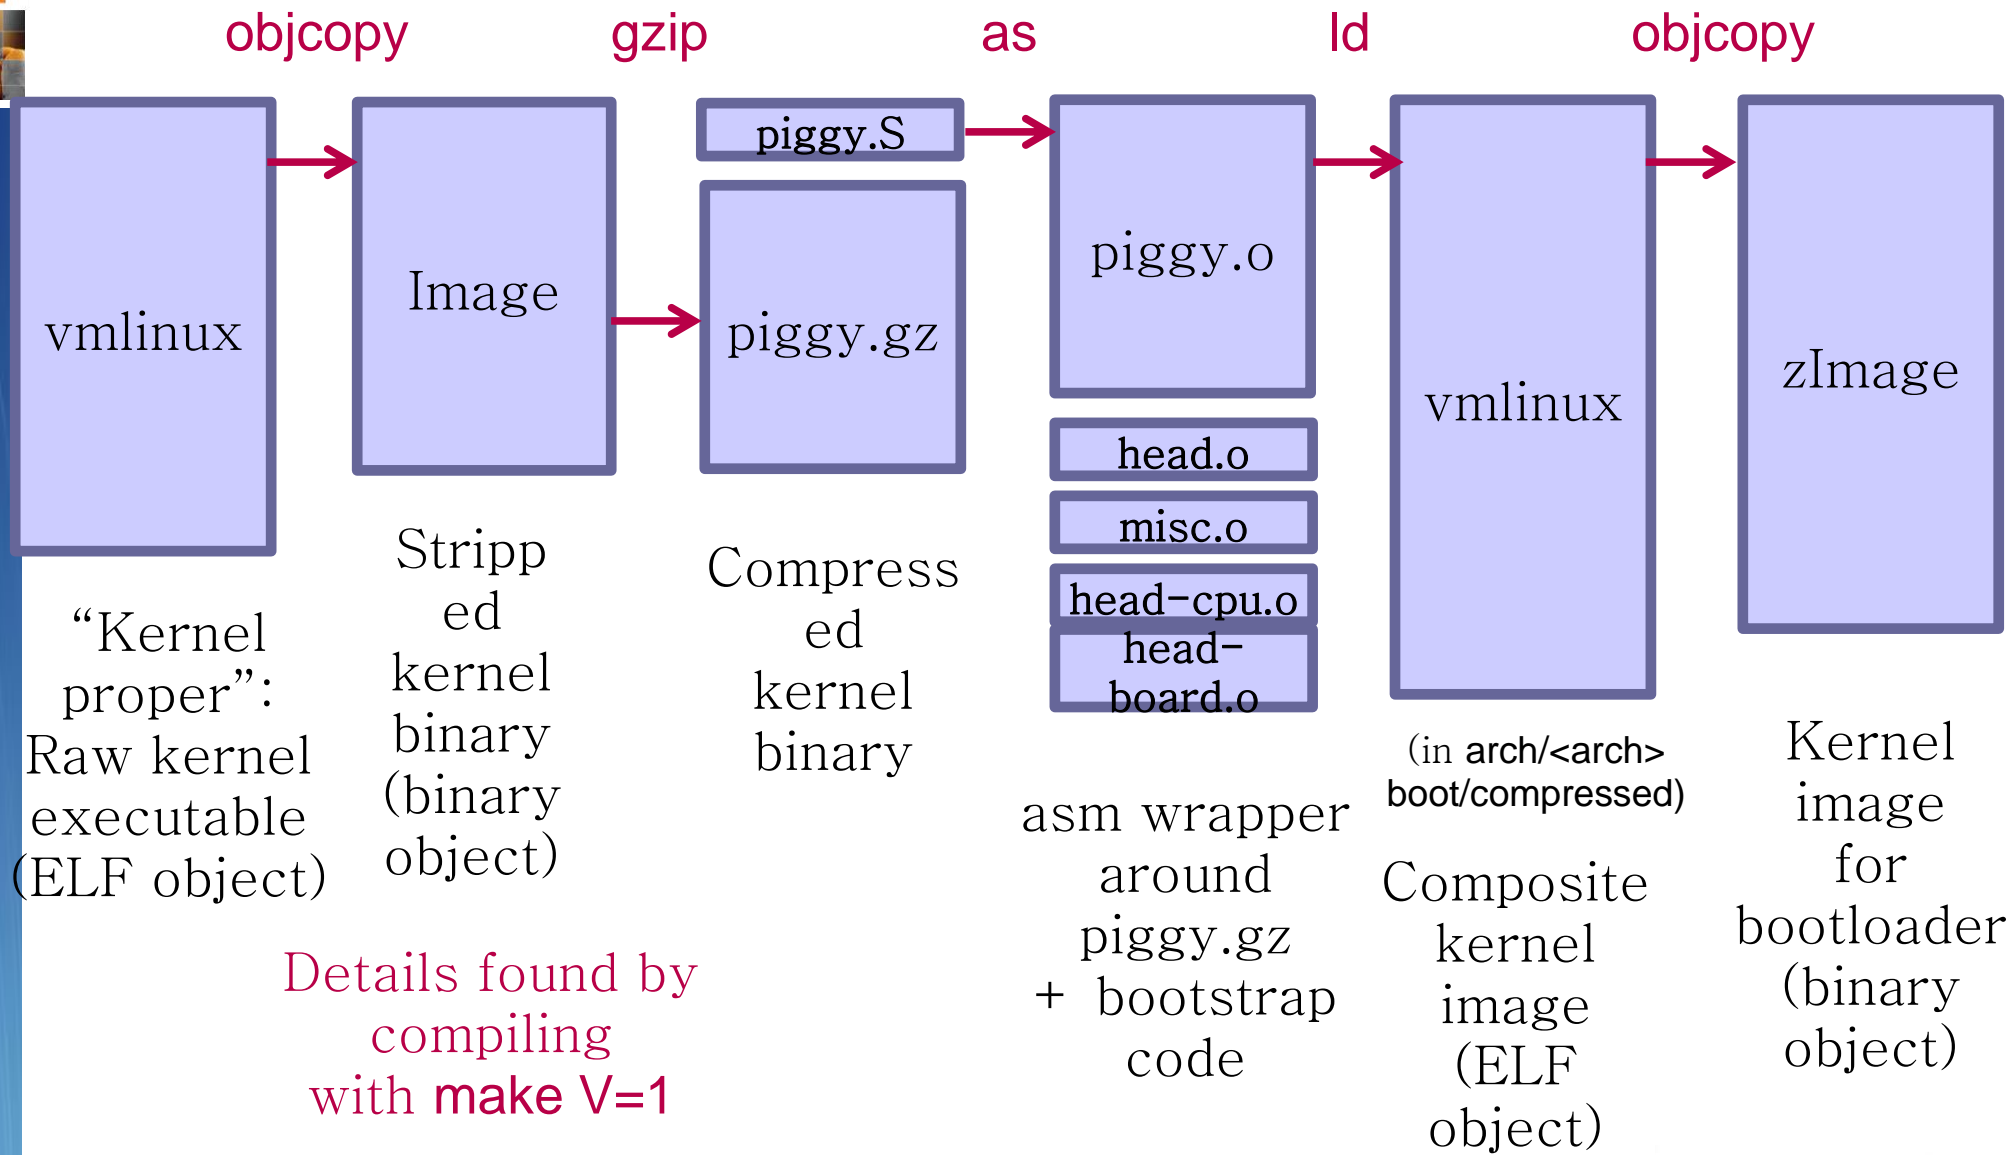
# Kernel Initialization



# From bootloader to userspace



# Kernel bootstrap



# Bootstrap code

## ▶ head.o:

Architecture specific initialization code.  
This is what is executed by the bootloader

## ▶ head-cpu.o

CPU specific initialization code

## ▶ head-board.o

Board specific initialization code

## ▶ misc.o:

Decompression routines

# Bootstrap code tasks

Main work done by [head.o](#):

- ▶ Check the architecture, processor and machine type.
- ▶ Configure the MMU, create page table entries and enable virtual memory.
- ▶ Calls the [start\\_kernel](#) function in [init/main.c](#).

Same code for all architectures.

Anybody interesting in kernel startup should study this file!

# start\_kernel main actions

- ▶ Calls `setup_arch(&command_line)` (function defined in `arch/<arch>/kernel/setup.c`), copying the command line from where the bootloader left it.
- ▶ On `arm`, this function calls `setup_processor` (in which CPU information is displayed) and `setup_machine` (locating the machine in the list of supported machines).
- ▶ Initializes the console as early as possible (to get error messages)
- ▶ Initializes many subsystems (see the code)
- ▶ Eventually calls `rest_init`.



# rest\_init: starting the init process

Starting a new kernel thread which will later become the init process

```
static void noline __init_refok rest_init(void)
{
    __releases(kernel_lock)

    int pid;

    kernel_thread(kernel_init, NULL, CLONE_FS | CLONE_SIGHAND);
    numa_default_policy();
    pid = kernel_thread(kthreadd, NULL, CLONE_FS | CLONE_FILES);
    kthreadd_task = find_task_by_pid(pid);
    unlock_kernel();

    /*
     * The boot idle thread must execute schedule()
     * at least one to get things moving:
     */
    preempt_enable_no_resched();
    schedule();
    preempt_disable();

    /* Call into cpu_idle with preempt disabled */
    cpu_idle();
}
```

Source: Linux  
2.6.22

# kernel\_init

kernel\_init does 2 main things:

► Call do\_basic\_setup

Now that kernel services are ready, start device initialization:

```
static void __init do_basic_setup(void)
{
    /* drivers will send hotplug events */
    init_workqueues();
    usermodehelper_init();
    driver_init();
    init_irq_proc();
    do_initcalls();
}
```

► Call init\_post





# do\_initcalls

Calls pluggable hooks registered with the macros below.  
Advantage: the generic code doesn't have to know about them.

```
/*
 * A "pure" initcall has no dependencies on anything else, and purely
 * initializes variables that couldn't be statically initialized.
 *
 * This only exists for built-in code, not for modules.
 */
#define pure_initcall(fn)      __define_initcall("0",fn,1)

#define core_initcall(fn)      __define_initcall("1",fn,1)
#define core_initcall_sync(fn) __define_initcall("1s",fn,1s)
#define postcore_initcall(fn) __define_initcall("2",fn,2)
#define postcore_initcall_sync(fn) __define_initcall("2s",fn,2s)
#define arch_initcall(fn)      __define_initcall("3",fn,3)
#define arch_initcall_sync(fn) __define_initcall("3s",fn,3s)
#define subsys_initcall(fn)     __define_initcall("4",fn,4)
#define subsys_initcall_sync(fn) __define_initcall("4s",fn,4s)
#define fs_initcall(fn)         __define_initcall("5",fn,5)
#define fs_initcall_sync(fn)    __define_initcall("5s",fn,5s)
#define rootfs_initcall(fn)     __define_initcall("rootfs",fn,rootfs)
#define device_initcall(fn)     __define_initcall("6",fn,6)
#define device_initcall_sync(fn) __define_initcall("6s",fn,6s)
#define late_initcall(fn)       __define_initcall("7",fn,7)
#define late_initcall_sync(fn)  __define_initcall("7s",fn,7s)
```



# initcall example

From arch/arm/mach-pxa/lpd270.c

```
static int __init lpd270_irq_device_init(void)
{
    int ret = sysdev_class_register(&lpd270_irq_sysclass);
    if (ret == 0)
        ret = sysdev_register(&lpd270_irq_device);
    return ret;
}

device_initcall(lpd270_irq_device_init);
```



# init\_post

The last step of Linux booting

- ▶ First tries to open a console
- ▶ Then tries to run the init process, effectively turning the current kernel thread into the userspace init process.

# init\_post code

```
static int noinline init_post(void)
{
    free_initmem();
    unlock_kernel();
    mark_rodata_ro();
    system_state = SYSTEM_RUNNING;
    numa_default_policy();

    if (sys_open((const char __user *) "/dev/console", O_RDWR, 0) < 0)
        printk(KERN_WARNING "Warning: unable to open an initial console.\n");

    (void) sys_dup(0);
    (void) sys_dup(0);

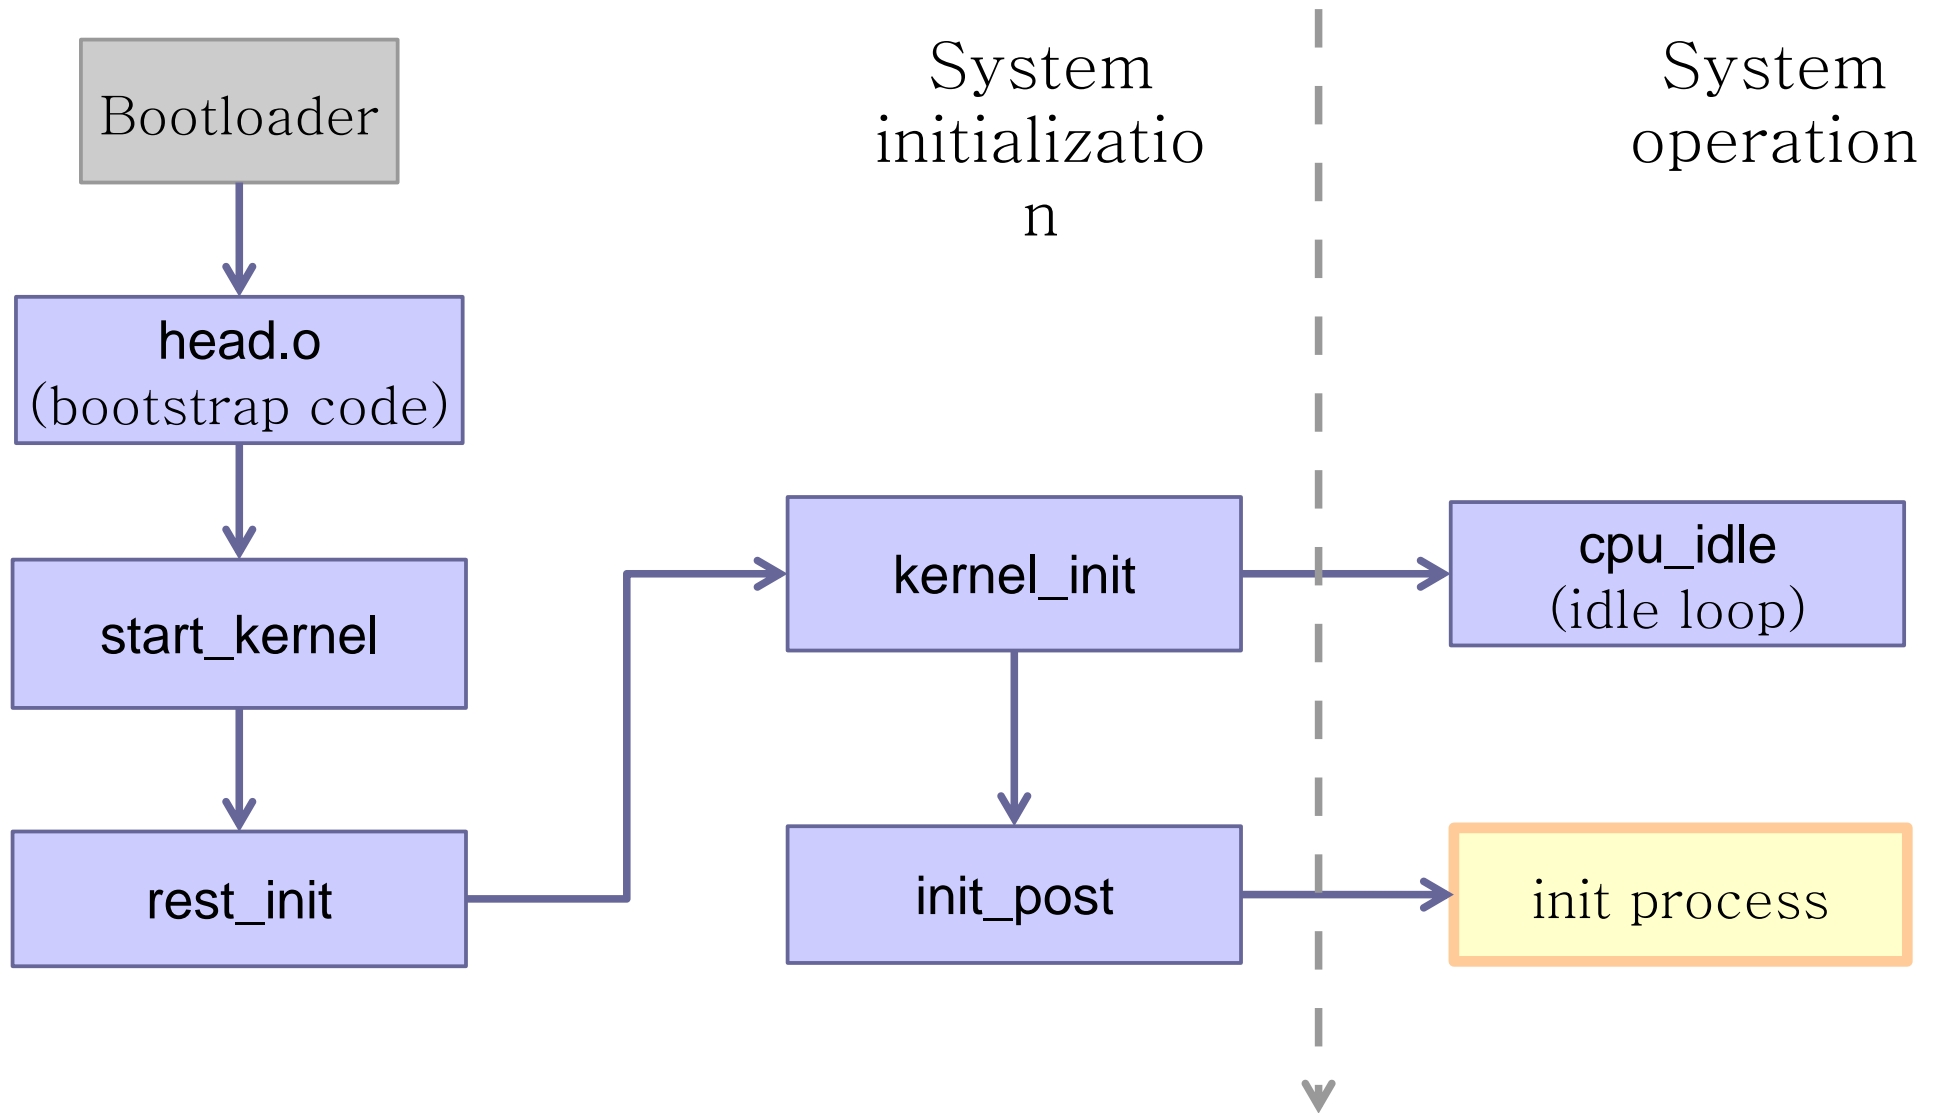
    if (ramdisk_execute_command) {
        run_init_process(ramdisk_execute_command);
        printk(KERN_WARNING "Failed to execute %s\n",
                ramdisk_execute_command);
    }

    /*
     * We try each of these until one succeeds.
     *
     * The Bourne shell can be used instead of init if we are
     * trying to recover a really broken machine.
     */
    if (execute_command) {
        run_init_process(execute_command);
        printk(KERN_WARNING "Failed to execute %s. Attempting "
                "defaults...\n", execute_command);
    }
    run_init_process("/sbin/init");
    run_init_process("/etc/init");
    run_init_process("/bin/init");
    run_init_process("/bin/sh");

    panic("No init found. Try passing init= option to kernel.");
}
```

Source:  
init/main.c  
in Linux 2.6.22

# Kernel initialization graph



# Kernel initialization - What to remember

- ▶ The bootloader executes bootstrap code.
- ▶ Bootstrap code initializes the processor and board, and uncompresses the kernel code to RAM, and calls the kernel's `start_kernel` function.
- ▶ Copies the command line from the bootloader.
- ▶ Identifies the processor and machine.
- ▶ Initializes the console.
- ▶ Initializes kernel services (memory allocation, scheduling, file cache...)
- ▶ Creates a new kernel thread (future init process) and continues in the idle loop.
- ▶ Initializes devices and execute initcalls.



# Thank You