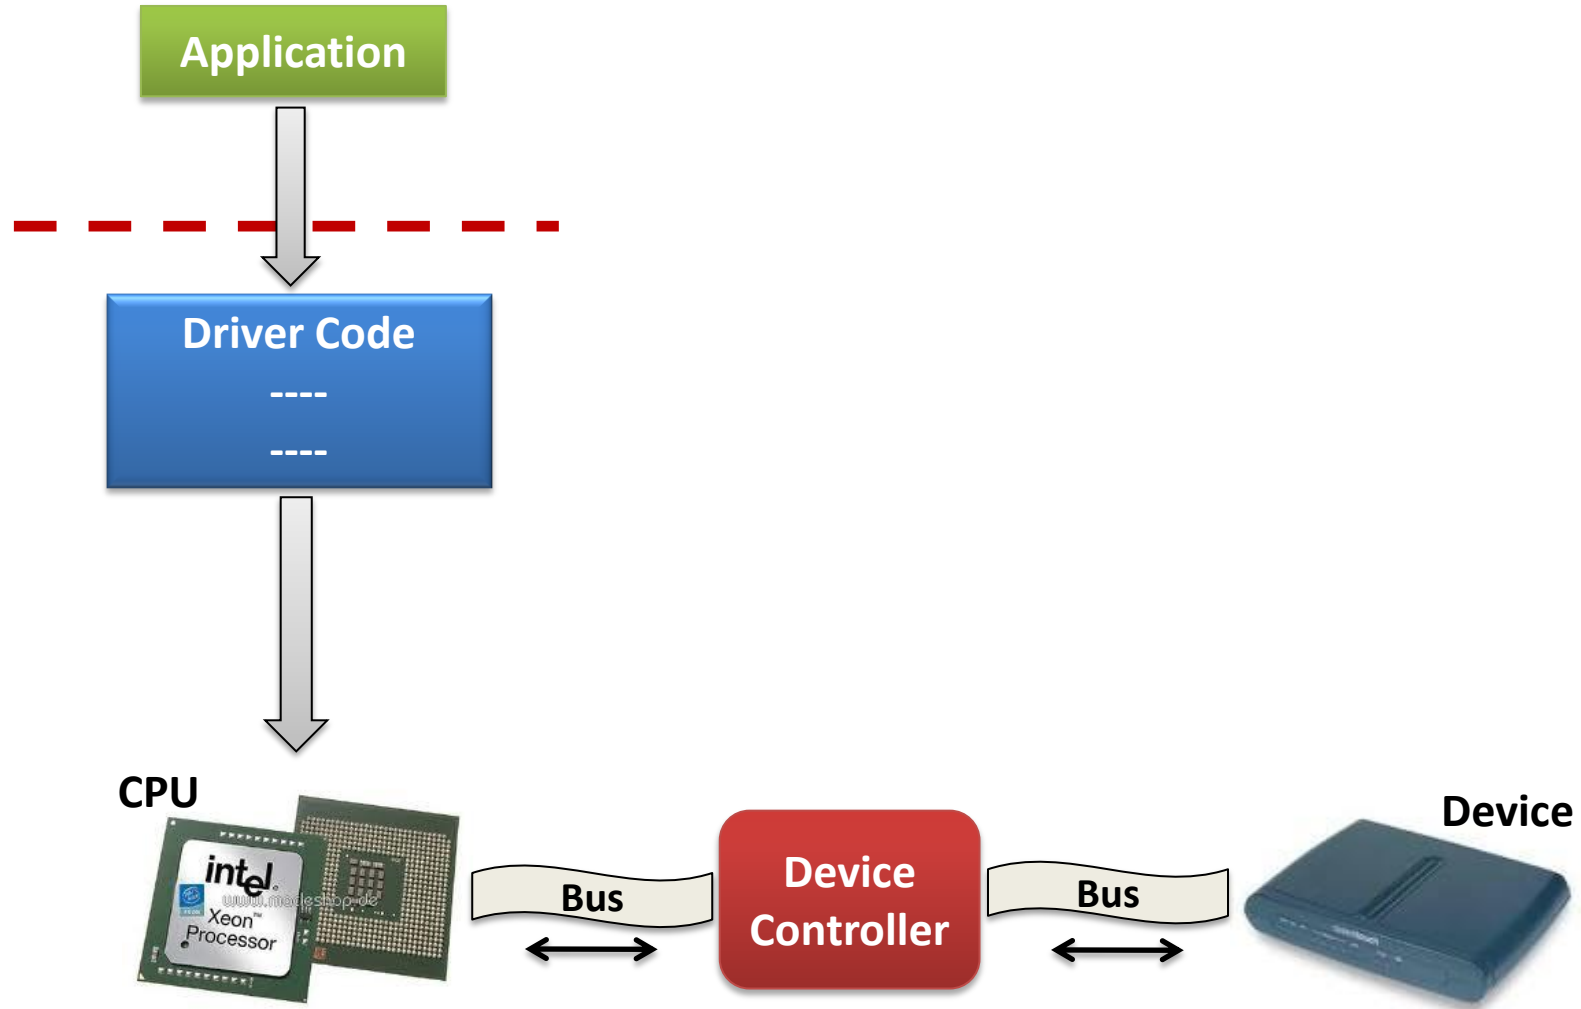


Device Drivers

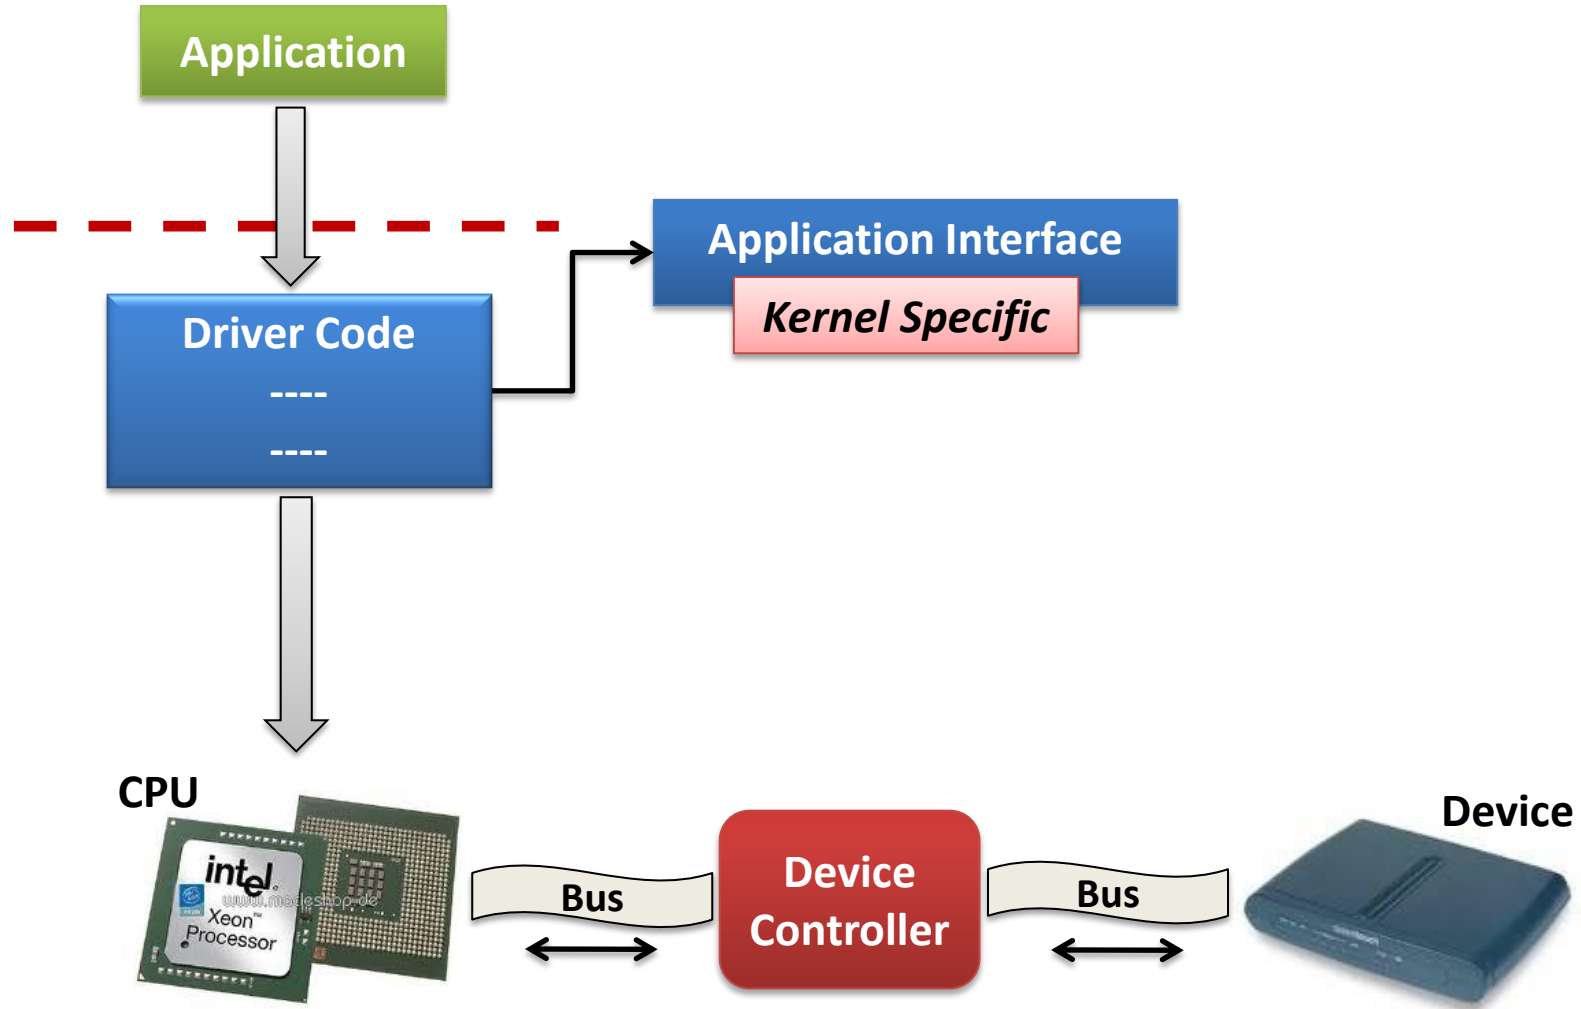
Device Driver Framework



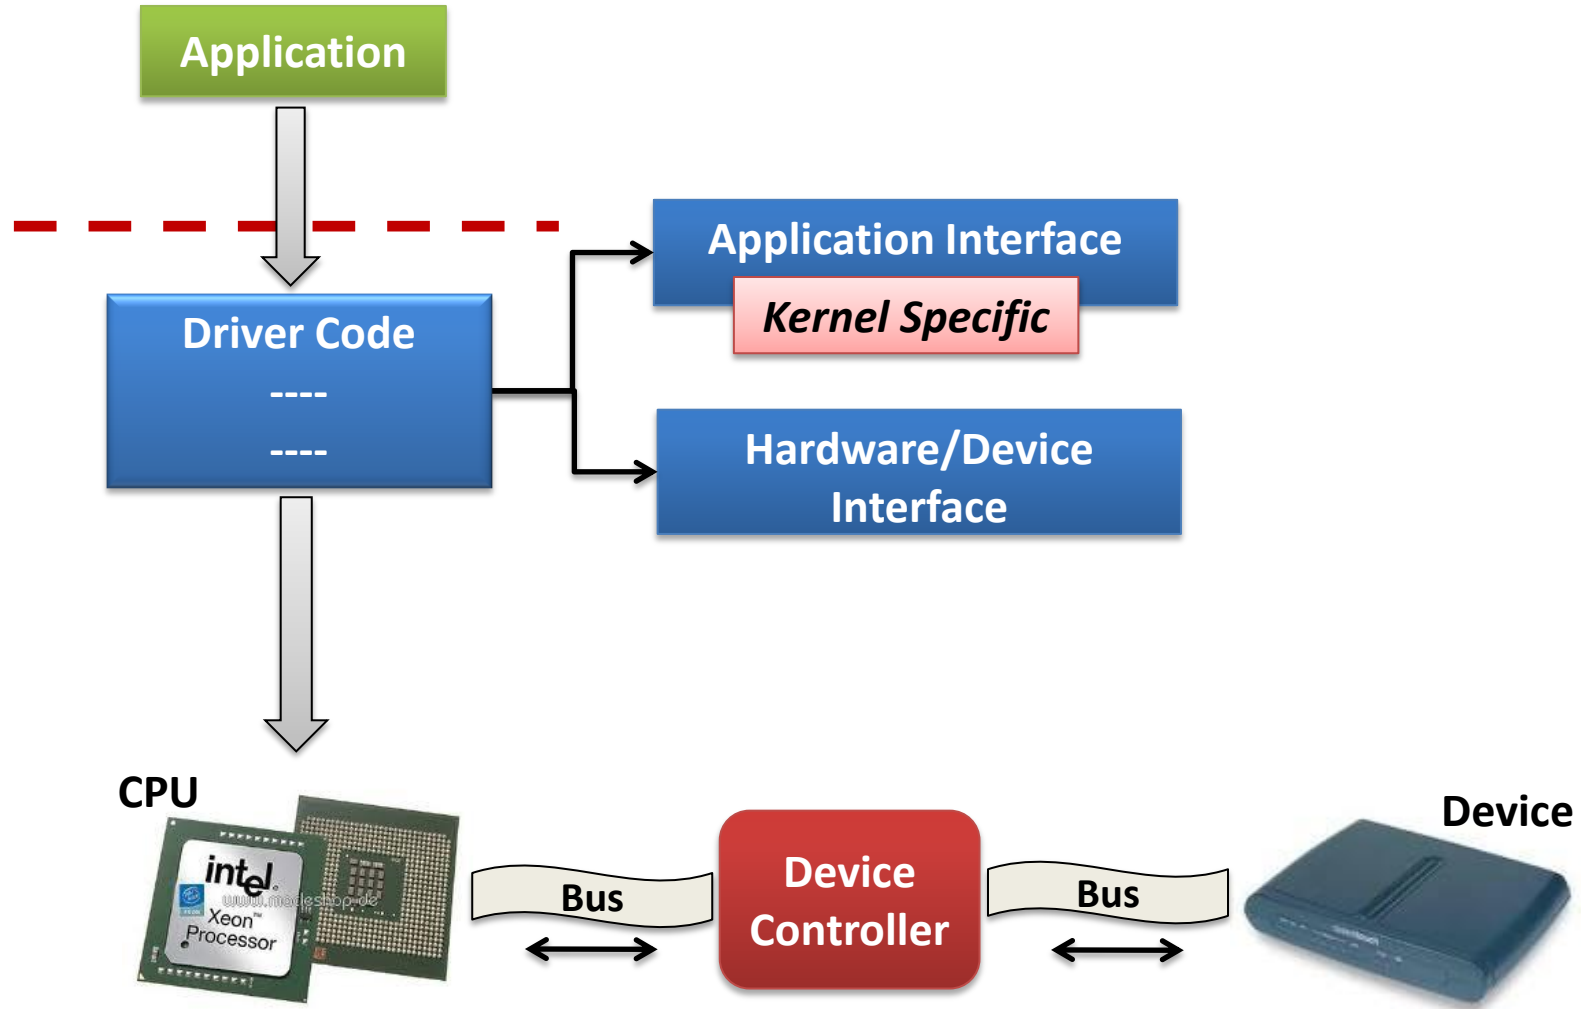
Device Driver Framework



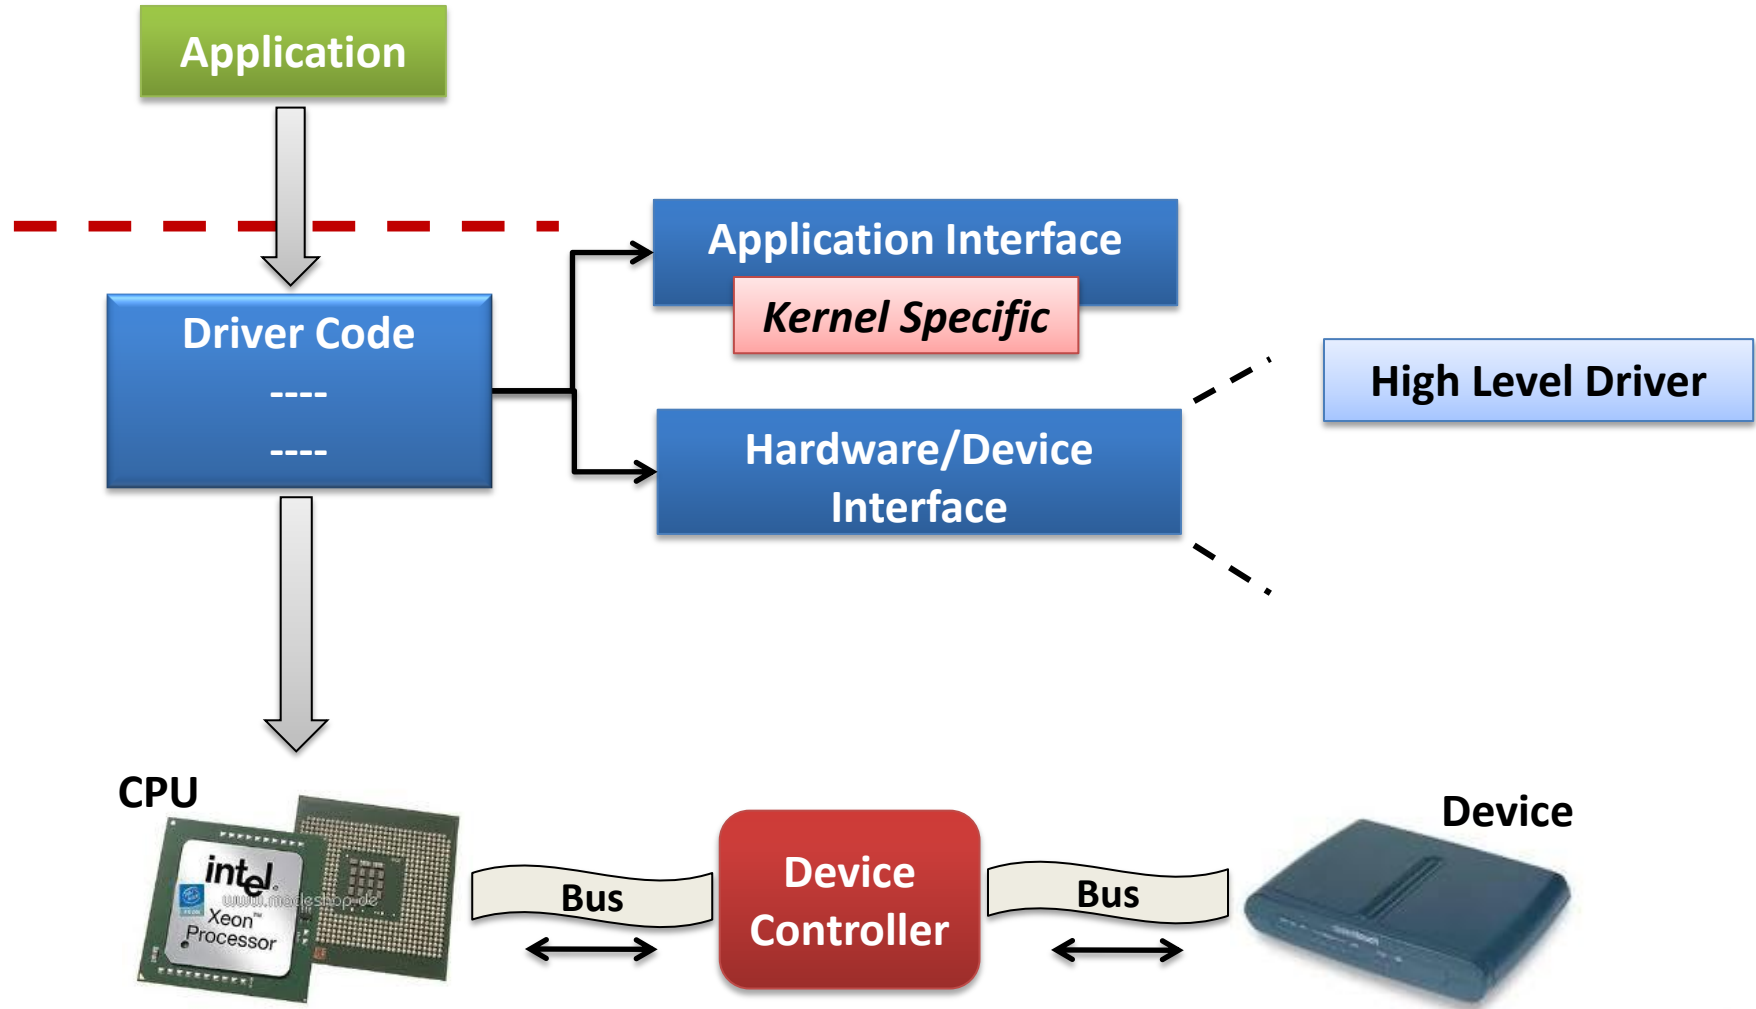
Device Driver Framework



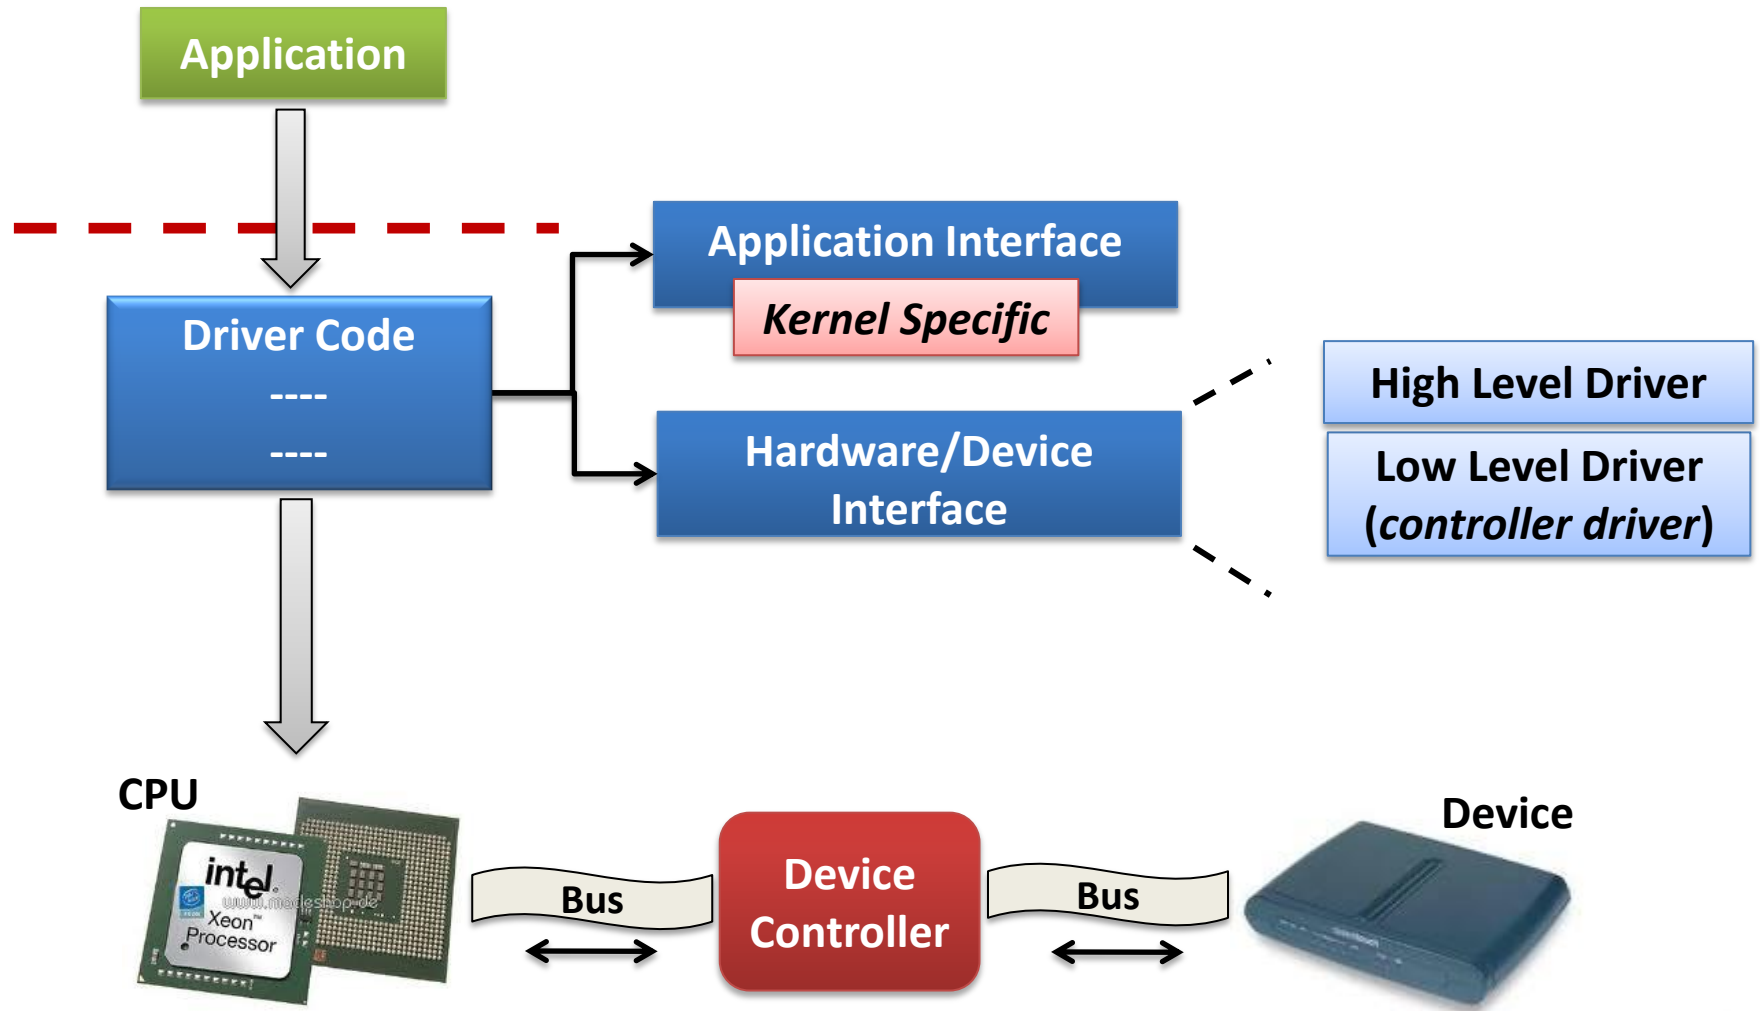
Device Driver Framework



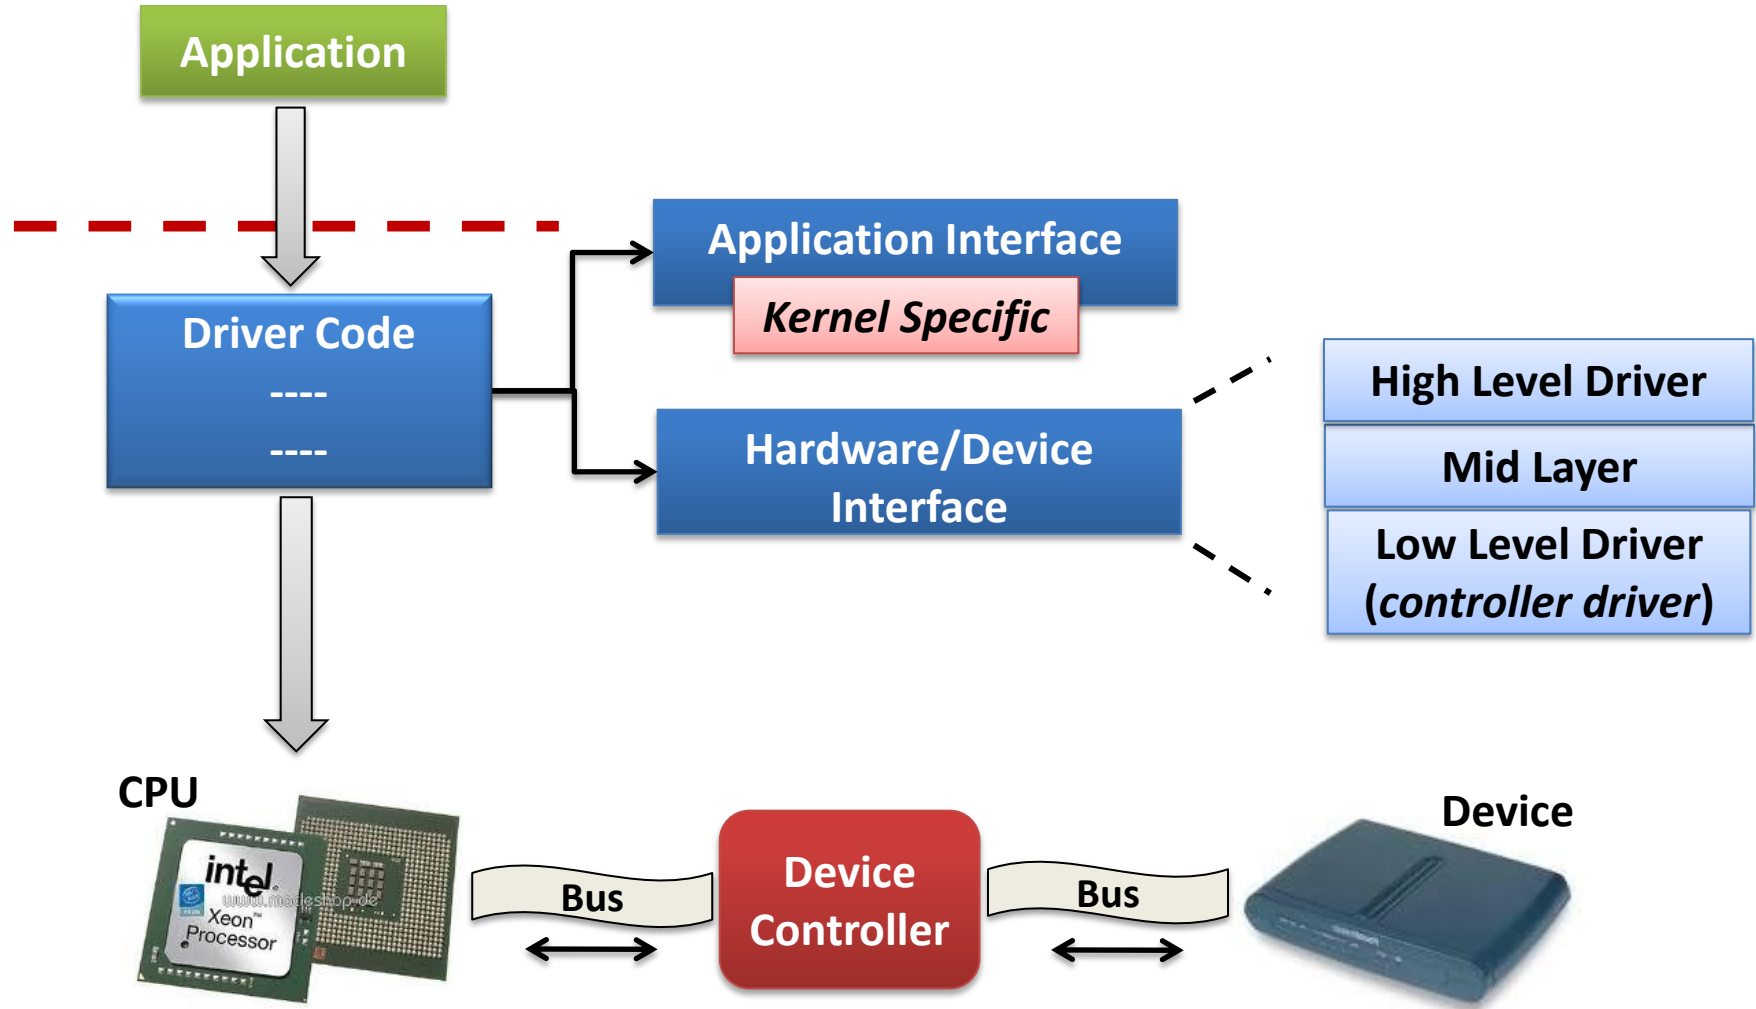
Device Driver Framework



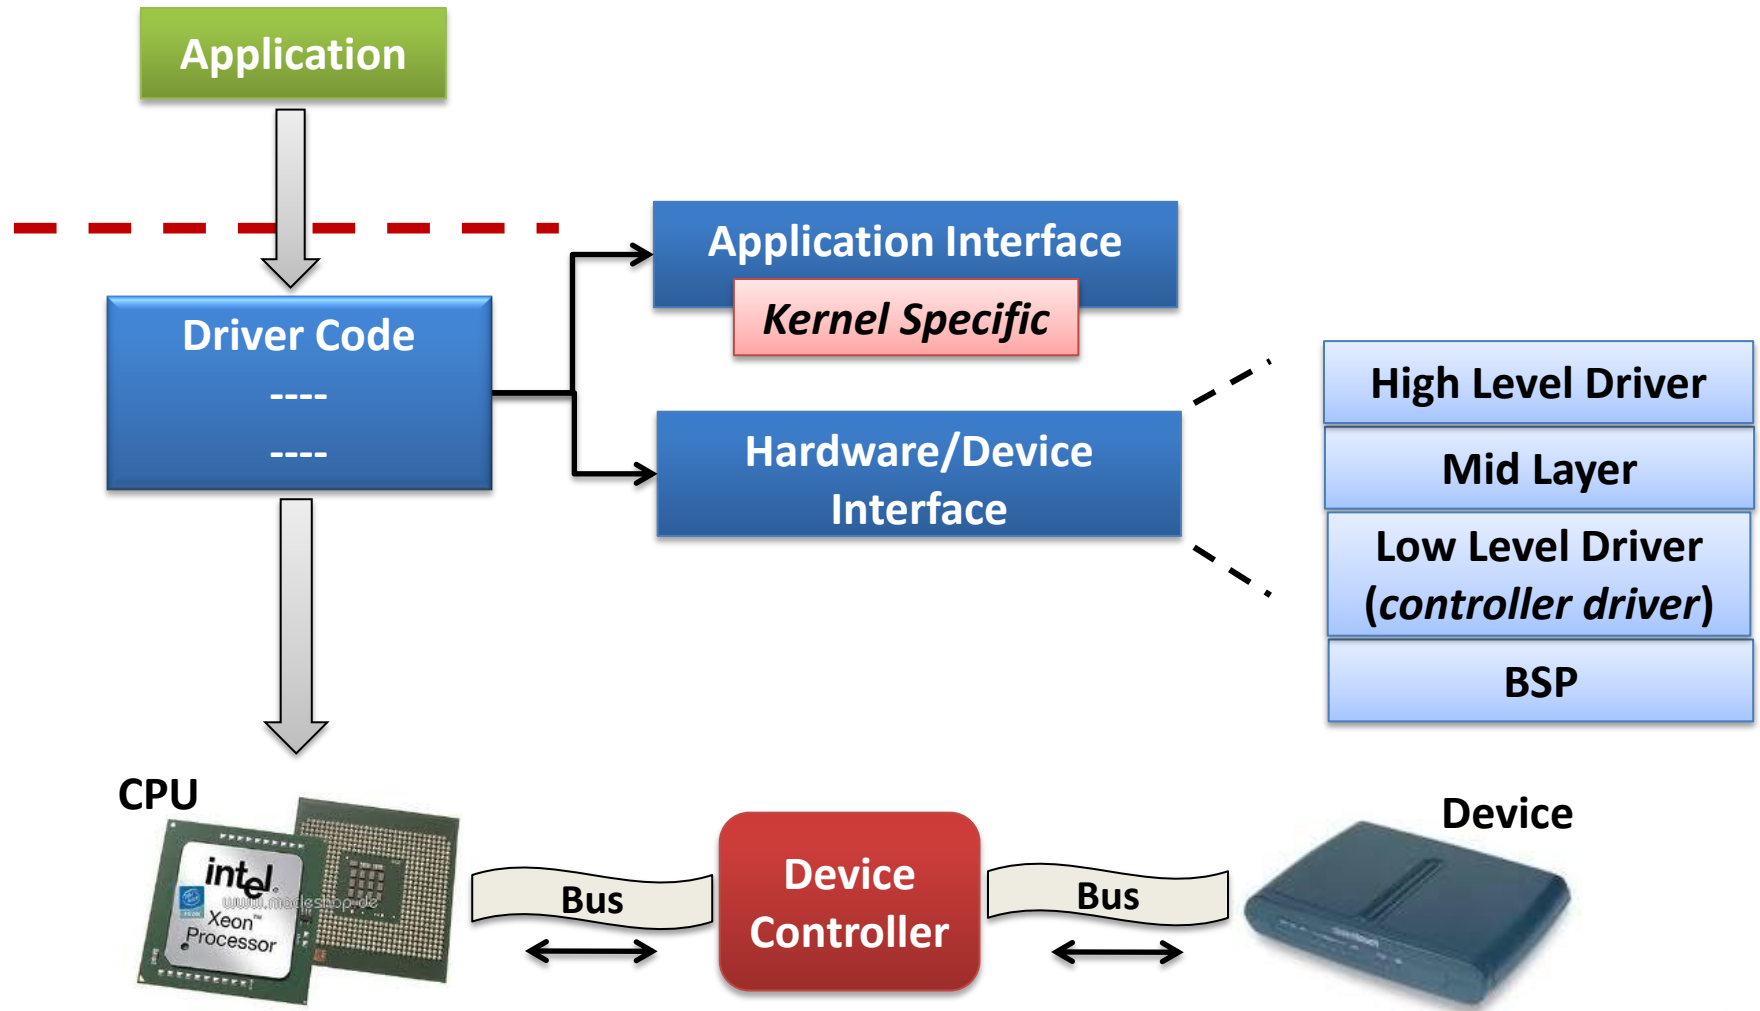
Device Driver Framework



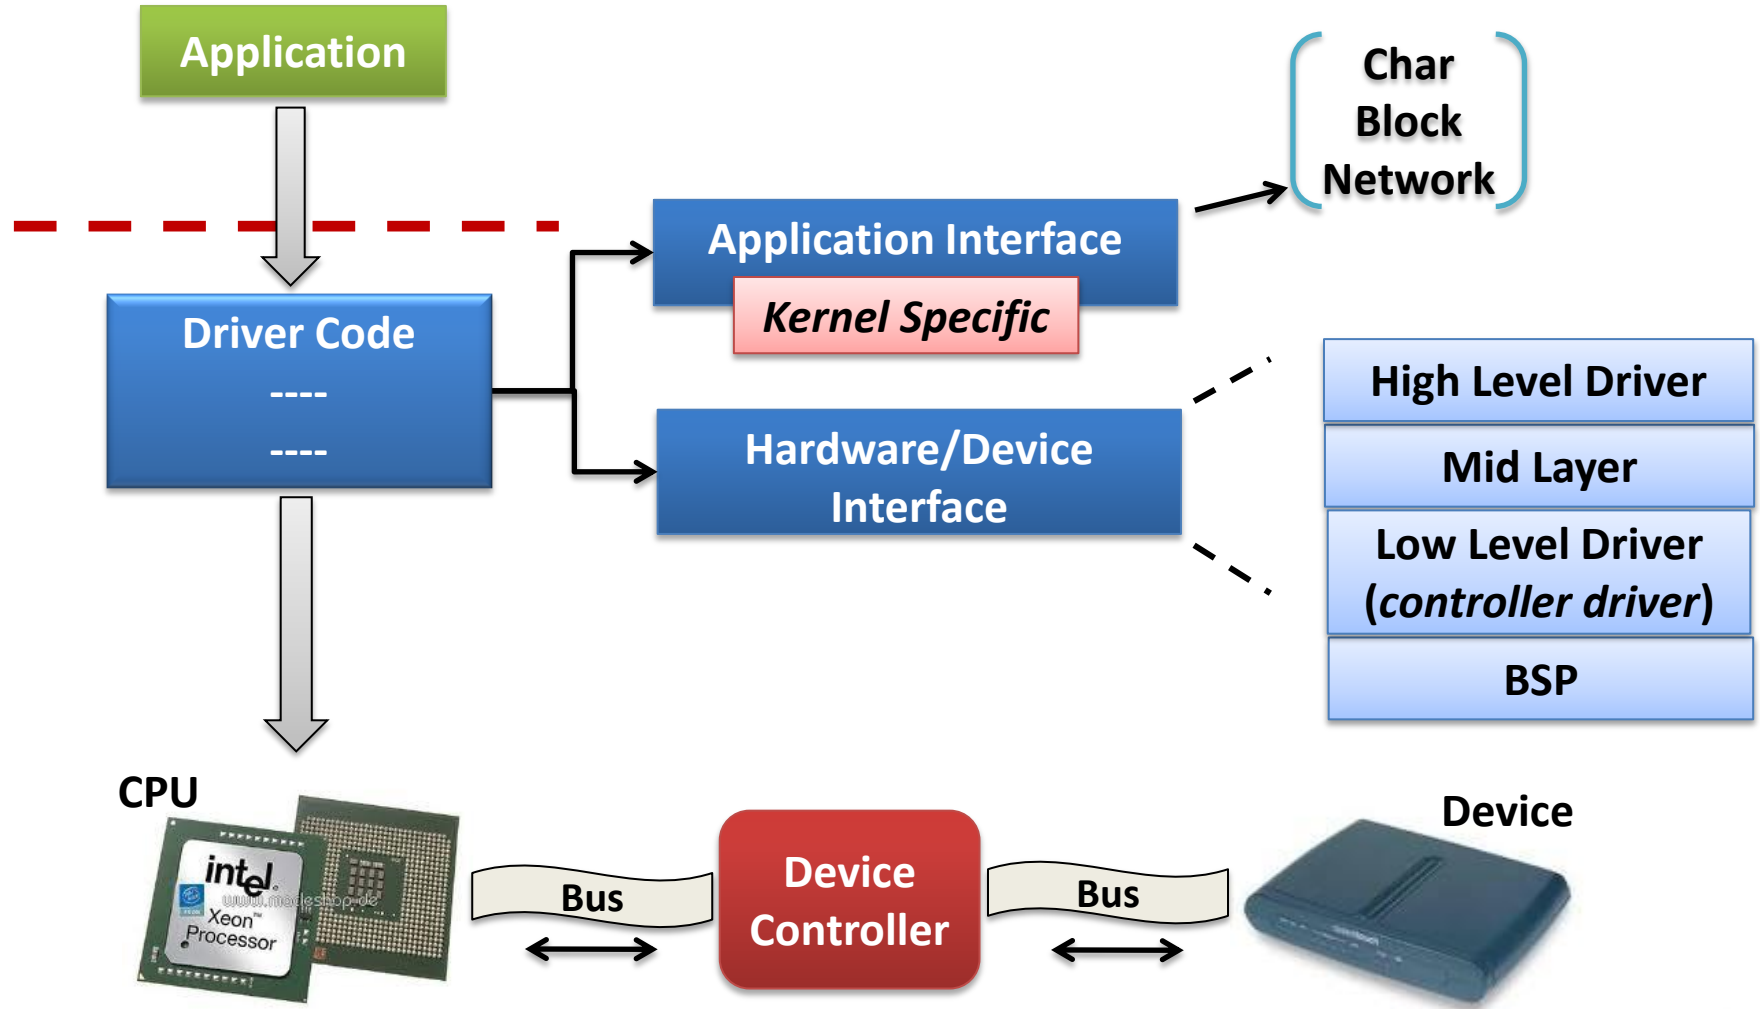
Device Driver Framework



Device Driver Framework



Device Driver Framework



Character Drivers

Character Driver

- ▶ Except for storage device drivers, most drivers for devices with input and output flows are implemented as character drivers.
- ▶ So, most drivers you will face will be character drivers
You will regret if you sleep during this part!

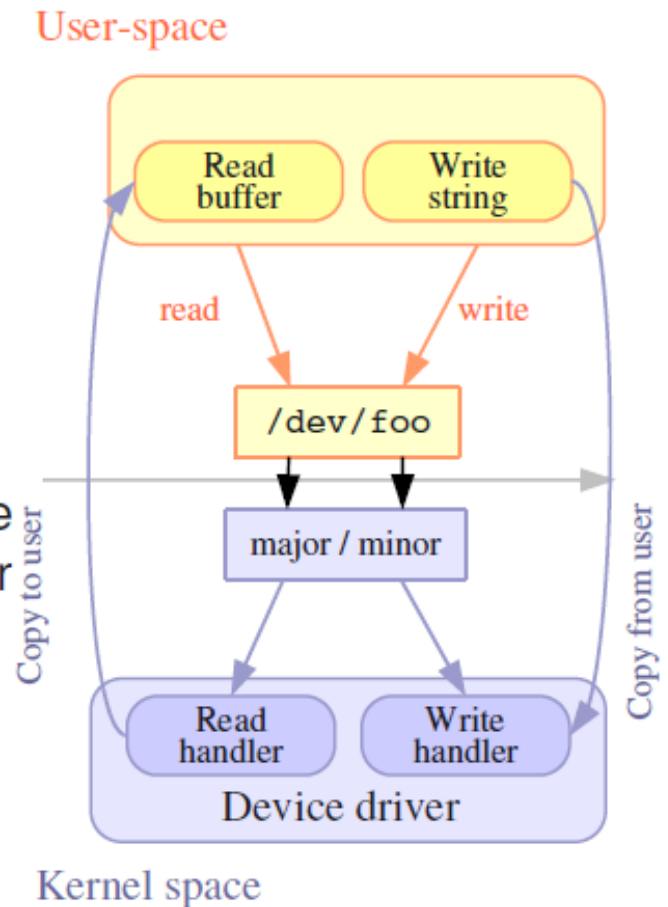
Creating a Character Driver

User-space needs

- ▶ The name of a device file in `/dev` to interact with the device driver through regular file operations (open, read, write, close...)

The kernel needs

- ▶ To know which driver is in charge of device files with a given major / minor number pair
- ▶ For a given driver, to have handlers (*“file operations”*) to execute when user-space opens, reads, writes or closes the device file.



Implementing Character Driver

- ▶ Four major steps
 - ▶ Implement operations corresponding to the system calls an application can apply to a file: file operations
 - ▶ Define a `file_operations` structure associating function pointers to their implementation in your driver
 - ▶ Reserve a set of major and minors for your driver
 - ▶ Tell the kernel to associate the reserved major and minor to your file operations
- ▶ This is a very common design scheme in the Linux kernel
 - ▶ A common kernel infrastructure defines a set of operations to be implemented by a driver and functions to register your driver
 - ▶ Your driver only needs to implement this set of well-defined operations

File Operations

- ▶ Before registering character devices, you have to define `file_operations` (called *fops*) for the device files.
- ▶ The `file_operations` structure is generic to all files handled by the Linux kernel. It contains many operations that aren't needed for character drivers.
- ▶ Here are the most important operations for a character driver. All of them are optional.

```
struct file_operations {
    [...]
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
    long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *);
    int (*release) (struct inode *, struct file *);
    [...]
};
```

open() and release()

▶ `int foo_open (struct inode *i, struct file *f)`

▶ Called when user-space opens the device file.

▶ `inode` is a structure that uniquely represent a file in the system (be it a regular file, a directory, a symbolic link, a character or block device)

▶ `file` is a structure created every time a file is opened. Several file structures can point to the same `inode` structure.

▶ Contains informations like the current position, the opening mode, etc.

▶ Has a `void *private_data` pointer that one can freely use.

▶ A pointer to the file structure is passed to all other operations

▶ `int foo_release(struct inode *i, struct file *f)`

▶ Called when user-space closes the file.

Exchanging data with user-space

- ▶ Kernel code isn't allowed to directly access user-space memory, using `memcpy` or direct pointer dereferencing
 - ▶ Doing so does not work on some architectures
 - ▶ If the address passed by the application was invalid, the application would segfault
- ▶ To keep the kernel code portable and have proper error handling, your driver must use special kernel functions to exchange data with user-space

Exchanging data with user-space

▶ A single value

▶ `get_user(v, p);`

The kernel variable `v` gets the value pointer by the user-space pointer `p`

▶ `put_user(v, p);`

The value pointed by the user-space pointer `p` is set to the contents of the kernel variable `v`.

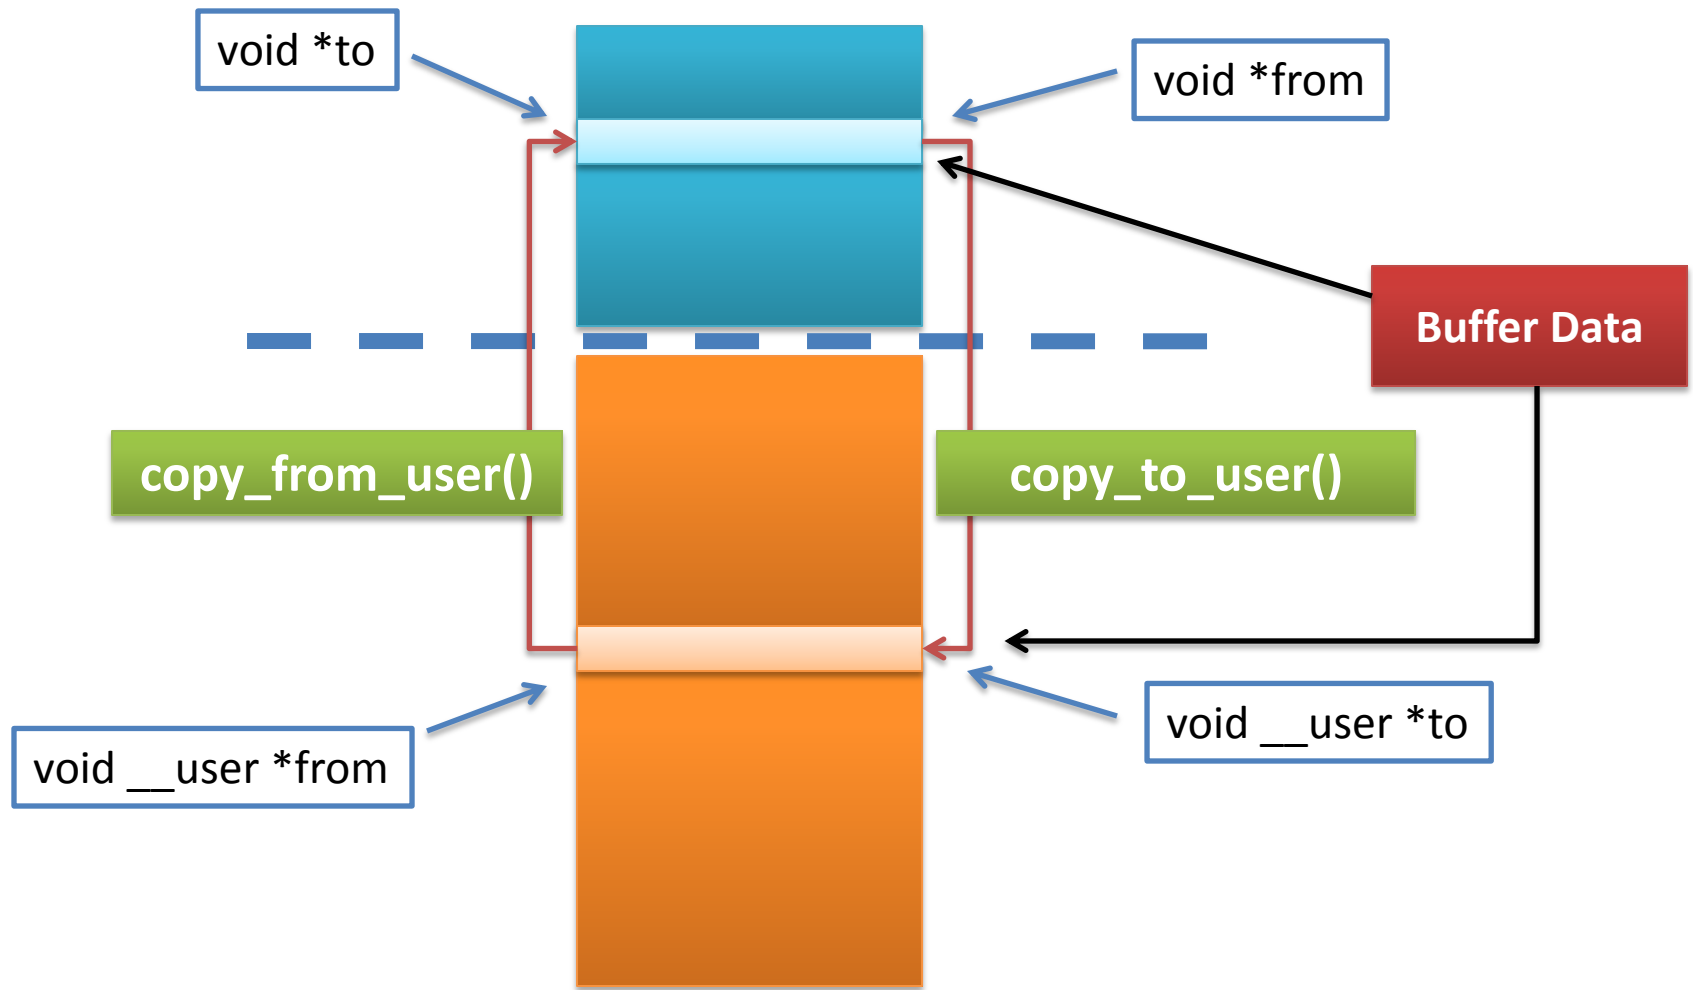
▶ A buffer

▶ `unsigned long copy_to_user(void __user *to,
const void *from, unsigned long n);`

▶ `unsigned long copy_from_user(void *to,
const void __user *from, unsigned long n);`

▶ The return value must be checked. Zero on success, non-zero on failure. If non-zero, the convention is to return `-EFAULT`.

Exchanging data with user-space



read operation example

```
static ssize_t
acme_read(struct file *file, char __user *buf, size_t count, loff_t *ppos)
{
    /* The acme_buf address corresponds to a device I/O memory area */
    /* of size acme_bufsize, obtained with ioremap() */
    int remaining_size, transfer_size;

    remaining_size = acme_bufsize - (int) (*ppos); // bytes left to transfer
    if (remaining_size == 0) { /* All read, returning 0 (End Of File) */
        return 0;
    }

    /* Size of this transfer */
    transfer_size = min(remaining_size, (int) count);

    if (copy_to_user(buf /* to */, acme_buf + *ppos /* from */, transfer_size)) {
        return -EFAULT;
    } else { /* Increase the position in the open file */
        *ppos += transfer_size;
        return transfer_size;
    }
}
```

Read method

write operation example

```
static ssize_t
acme_write(struct file *file, const char __user *buf, size_t count, loff_t *ppos)
{
    int remaining_bytes;

    /* Number of bytes not written yet in the device */
    remaining_bytes = acme_bufsize - (*ppos);

    if (count > remaining_bytes) {
        /* Can't write beyond the end of the device */
        return -EIO;
    }

    if (copy_from_user(acme_buf + *ppos /* to */, buf /* from */, count)) {
        return -EFAULT;
    } else {
        /* Increase the position in the open file */
        *ppos += count;
        return count;
    }
}
```

write method

unlocked_ioctl()

```
long unlocked_ioctl(struct file *f,
                    unsigned int cmd, unsigned long arg)
```

- ▶ Associated to the `ioctl()` system call
Called `unlocked` because it doesn't hold the Big Kernel Lock.
- ▶ Allows to extend the driver capabilities beyond the limited read/write API
- ▶ For example: changing the speed of a serial port, setting video output format, querying a device serial number...
- ▶ `cmd` is a number identifying the operation to perform
- ▶ `arg` is the optional argument passed as third argument of the `ioctl()` system call. Can be an integer, an address, etc.
- ▶ The semantic of `cmd` and `arg` is driver-specific.

ioctl() example - Kernel side

```
static long phantom_ioctl(struct file *file, unsigned int cmd,
                          unsigned long arg)
{
    struct phm_reg r;
    void __user *argp = (void __user *)arg;

    switch (cmd) {
    case PHN_SET_REG:
        if (copy_from_user(&r, argp, sizeof(r)))
            return -EFAULT;
        /* Do something */
        break;
    case PHN_GET_REG:
        if (copy_to_user(argp, &r, sizeof(r)))
            return -EFAULT;

        /* Do something */
        break;
    default:
        return -ENOTTY;
    }

    return 0;
}
```

ioctl() example - application side

```
int main(void)
{
    int fd, ret;
    struct phm_reg reg;

    fd = open("/dev/phantom");
    assert(fd > 0);

    reg.field1 = 42;
    reg.field2 = 67;

    ret = ioctl(fd, PHN_SET_REG, & reg);
    assert(ret == 0);

    return 0;
}
```


File operations definition example

Defining a `file_operations` structure:

```
#include <linux/fs.h>

static struct file_operations acme_fops =
{
    .owner = THIS_MODULE,
    .read = acme_read,
    .write = acme_write,
};
```

You just need to supply the functions you implemented! Defaults for other functions (such as `open`, `release...`) are fine if you do not implement anything special.

dev_t data type

Kernel data type to represent a major / minor number pair

- ▶ Also called a *device number*.
- ▶ Defined in `<linux/kdev_t.h>`
Linux 2.6: 32 bit size (major: 12 bits, minor: 20 bits)
- ▶ Macro to compose the device number:
`MKDEV(int major, int minor);`
- ▶ Macro to extract the minor and major numbers:
`MAJOR(dev_t dev);`
`MINOR(dev_t dev);`

Registering device numbers

```
#include <linux/fs.h>

int register_chrdev_region(
    dev_t from,           /* Starting device number */
    unsigned count,       /* Number of device numbers */
    const char *name);    /* Registered name */
```

Returns 0 if the allocation was successful.

Example

```
static dev_t acme_dev = MKDEV(202, 128);

if (register_chrdev_region(acme_dev, acme_count, "acme")) {
    printk(KERN_ERR "Failed to allocate device number\n");
    ...
}
```



Registering device numbers

If you don't have fixed device numbers assigned to your driver

- ▶ Better not to choose arbitrary ones.
There could be conflicts with other drivers.
- ▶ The kernel API offers a `alloc_chrdev_region` function to have the kernel allocate free ones for you. You can find the allocated major number in `/proc/devices`.

Information on registered numbers

Registered devices are visible in `/proc/devices`:

| Character devices: | Block devices: |
|--------------------|----------------|
| 1 mem | 1 ramdisk |
| 4 /dev/vc/0 | 3 ide0 |
| 4 tty | 8 sd |
| 4 ttyS | 9 md |
| 5 /dev/tty | 22 ide1 |
| 5 /dev/console | 65 sd |
| 5 /dev/ptmx | 66 sd |
| 6 lp | 67 sd |
| 10 misc | 68 sd |
| 13 input | |
| 14 sound | |
| ... | |

Major
number
Registered
name

Character Driver Registration

- ▶ The kernel represents character drivers with a `cdev` structure
- ▶ Declare this structure globally (within your module):

```
#include <linux/cdev.h>
static struct cdev acme_cdev;
```
- ▶ In the init function, initialize the structure:

```
cdev_init(&acme_cdev, &acme_fops);
```

Character Driver Registration

- ▶ Then, now that your structure is ready, add it to the system:

```
int cdev_add(
    struct cdev *p,      /* Character device structure */
    dev_t dev,          /* Starting device major / minor number */
    unsigned count);    /* Number of devices */
```

- ▶ After this function call, the kernel knows the association between the major/minor numbers and the file operations. Your device is ready to be used!

- ▶ Example (continued):

```
if (cdev_add(&acme_cdev, acme_dev, acme_count)) {
    printk (KERN_ERR "Char driver registration failed\n");
    ...
}
```

Character Device Un-registration

- ▶ First delete your character device:

```
void cdev_del(struct cdev *p);
```

- ▶ Then, and only then, free the device number:

```
void unregister_chrdev_region(dev_t from,  
unsigned count);
```

- ▶ Example (continued):

```
cdev_del(&acme_cdev);  
unregister_chrdev_region(acme_dev, acme_count);
```


Linux error codes

- ▶ The kernel convention for error management is
 - ▶ Return 0 on success
`return 0;`
 - ▶ Return a negative error code on failure
`return -EFAULT;`
- ▶ Error codes
 - ▶ `include/asm-generic/errno-base.h`
 - ▶ `include/asm-generic/errno.h`

Char driver example summary

```
static void *acme_buf;
static int acme_bufsize=8192;

static int acme_count=1;
static dev_t acme_dev = MKDEV(202,128);

static struct cdev acme_cdev;

static ssize_t acme_write(...) {...}

static ssize_t acme_read(...) {...}

static struct file_operations acme_fops =
{
    .owner = THIS_MODULE,
    .read = acme_read,
    .write = acme_write
};
```

Char driver example summary

Shows how to handle errors and deallocate resources in the right order!

```
static int __init acme_init(void)
{
    int err;
    acme_buf = ioremap (ACME_PHYS,
                       acme_bufsize);

    if (!acme_buf) {
        err = -ENOMEM;
        goto err_exit;
    }

    if (register_chrdev_region(acme_dev,
                              acme_count, "acme")) {
        err=-ENODEV;
        goto err_free_buf;
    }

    cdev_init(&acme_cdev, &acme_fops);

    if (cdev_add(&acme_cdev, acme_dev,
                acme_count)) {
        err=-ENODEV;
        goto err_dev_unregister;
    }
}
```

```
    return 0;

err_dev_unregister:
    unregister_chrdev_region(
        acme_dev, acme_count);
err_free_buf:
    iounmap(acme_buf);
err_exit:
    return err;
}

static void __exit acme_exit(void)
{
    cdev_del(&acme_cdev);
    unregister_chrdev_region(acme_dev,
                             acme_count);
    iounmap(acme_buf);
}
```

Char driver summary

Character driver writer

- Define the file operations callbacks for the device file: `read`, `write`, `ioctl`...
- In the module init function, reserve major and minor numbers with `register_chrdev_region()`, init a `cdev` structure with your file operations and add it to the system with `cdev_add()`.
- In the module exit function, call `cdev_del()` and `unregister_chrdev_region()`

Kernel

System administration

- Load the character driver module
 - Create device files with matching major and minor numbers if needed
- The device file is ready to use!

User-space

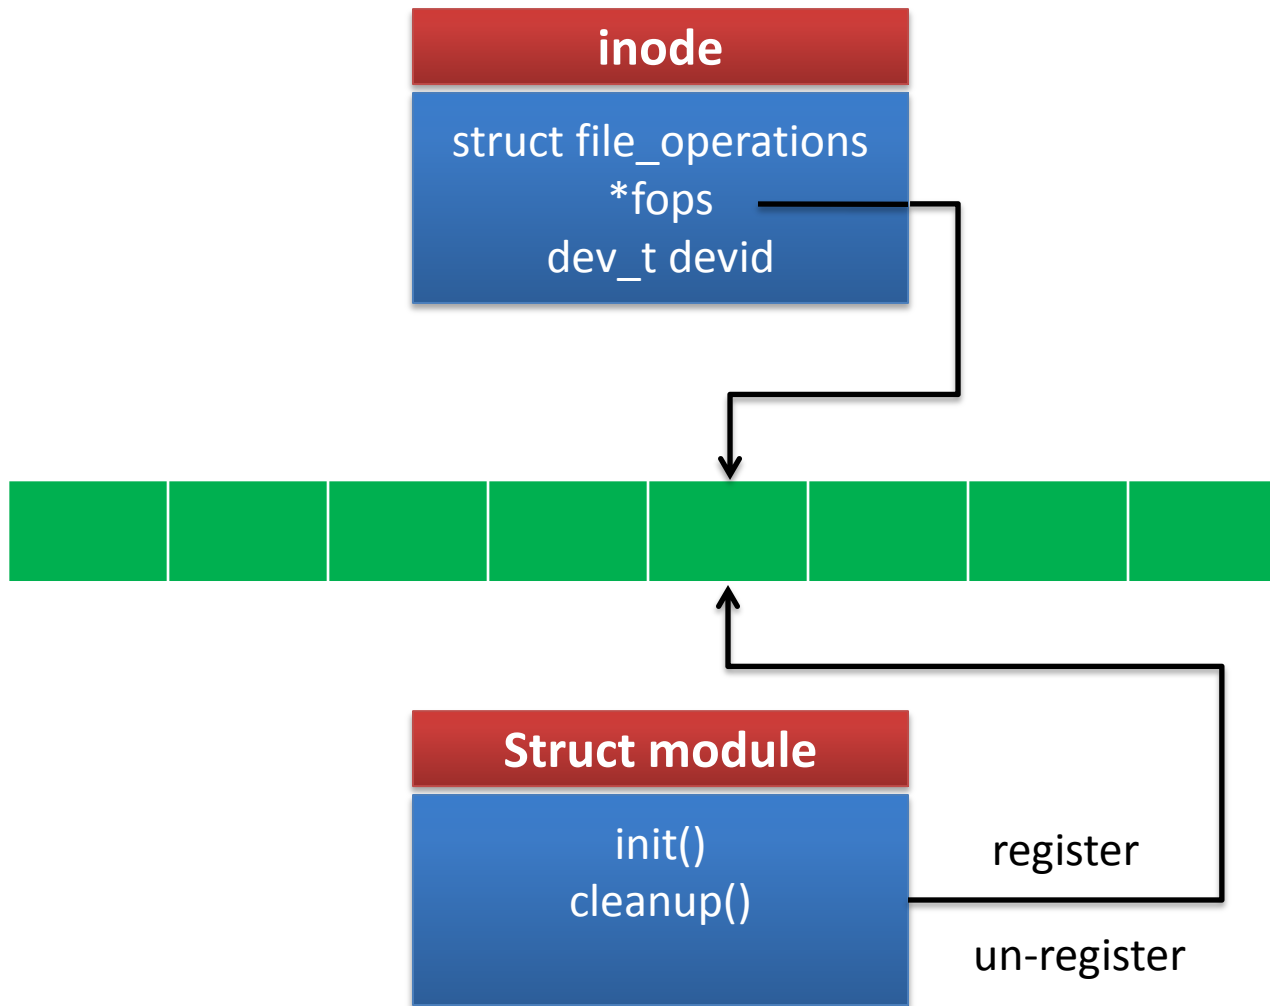
System user


- Open the device file, read, write, or send `ioctl`'s to it.

Kernel

- Executes the corresponding file operations

Kernel





Thank You