

Interrupts

Interrupts

- ✓ An interrupt is an event that **alters** the sequence of instructions executed by a processor in corresponding to electrical signals generated by HW circuits both **inside** and **outside** CPU
- ✓ **Interrupts**: asynchronous interrupts
Generated by HW devices (e.g., internal timers and I/O devices) at arbitrary times
- ✓ **Exceptions**: synchronous interrupts
Produced by CPU control unit only after completion of an executing instruction
E.g., divide-by-0, page faults

Exceptions

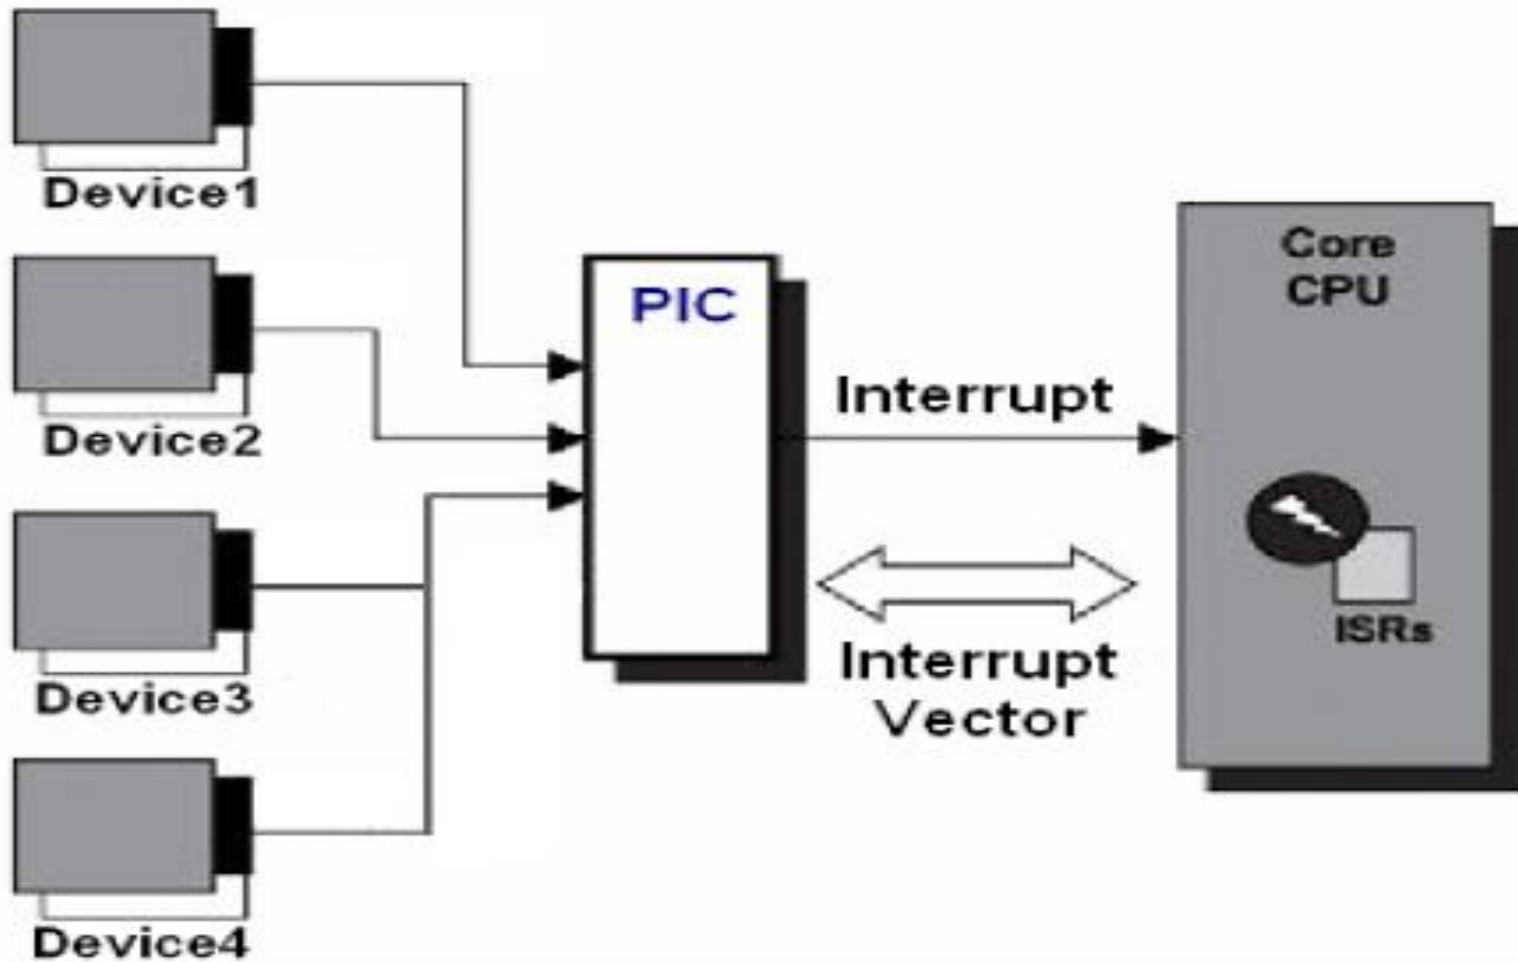
Processor-detected exceptions: when CPU detects anomalous condition while executing an instruction

- ✓ **Faults**: instructions causing faults
- ✓ **Traps** :Main usage: debugging purpose (e.g. reaching a breakpoint)
- ✓ **Aborts**: a serious error that may be unable to determine exact inst causing this error → terminate affected process

Programmed exceptions: occur at the request of programmer

- ✓ Triggered by **int**, **int3**, **into**, **bound** instructions
- ✓ Handled by control unit as traps
- ✓ Often called **SW interrupts**
- ✓ Usage: to implement system calls and to notify a debugger of a specific event

IO interrupts



IRQ'S

- ✓ Each HW device controller capable of issuing interrupts has an **output line IRQ**
- ✓ All existing IRQ lines are connected to the input pins of the **Interrupt Controller**
- ✓ Interrupt Controller (IC) executes


Monitoring IRQ lines, checking for raised signals

If a raised signal is detected on an IRQ line

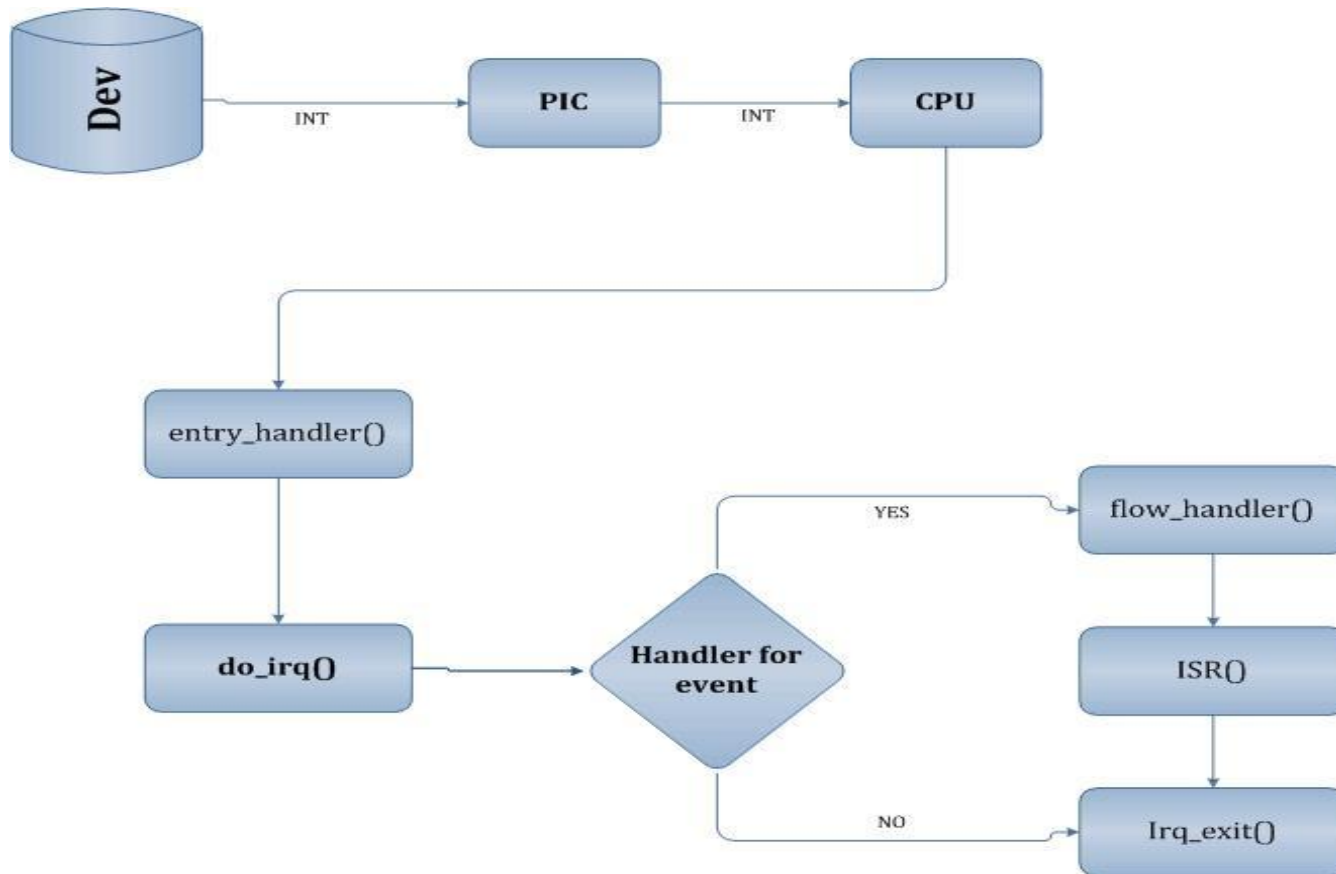
- ✓ Converts signal into a **corresponding vector**
- ✓ Stores vector in an IC I/O port, for CPU to read
- ✓ Sends a signal to CPU's INTR pin (i.e., issues an interrupt)
- ✓ CPU recognizes and writes to one of Programmable Interrupt Controller (PIC) I/O ports
- ✓ Clear INTR line

Go back to monitoring step

Linux interrupt Code



Interrupt Code Flow



```

struct irq_desc {
    unsigned int status;           /* IRQ line status */
    irq_flow_handler_t handle_irq; /* flow handler routine */
    struct irq_chip *chip;        /* chip operations */
    struct irqaction *action;     /* IRQ action ISR list */
    unsigned int depth;          /* nested irq disables */
    const char *name             /* name of flow control handler */
} ____cacheline_aligned irq_desc_t;

extern irq_desc_t irq_desc [NR_IRQS];    // global variable

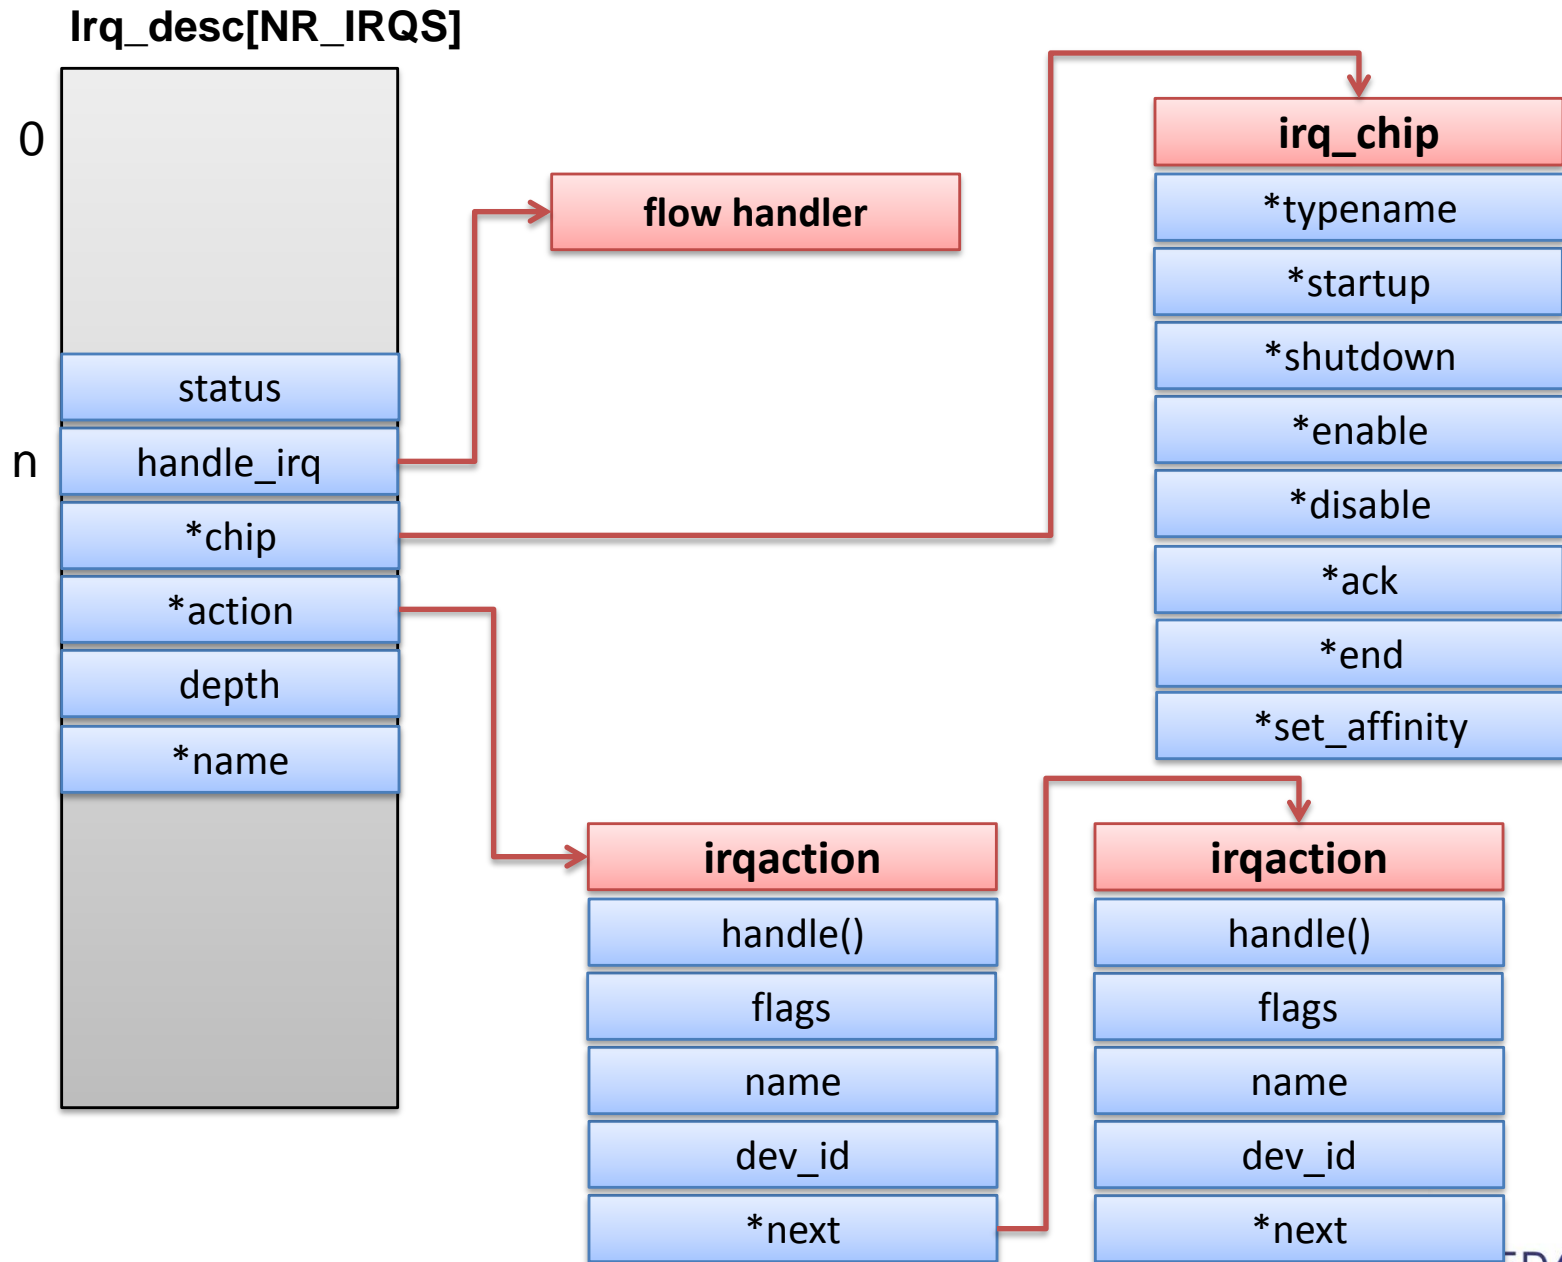
```

```

struct irq_chip {
    const char * typename;
    unsigned int (*startup) (unsigned int irq);
    void (*shutdown) (unsigned int irq);
    void (*enable) (unsigned int irq);
    void (*disable) (unsigned int irq);
    void (*ack) (unsigned int irq);
    void (*end) (unsigned int irq);
    void (*set_affinity) (unsigned int irq, cpumask_t dest);
};

```


IRQ Descriptors



IRQ Status Listing

```

/*
 * IRQ line status.
 */

#define IRQ_INPROGRESS      1      /* IRQ handler active - do not enter! */
#define IRQ_DISABLED        2      /* IRQ disabled - do not enter! */
#define IRQ_PENDING         4      /* IRQ pending - replay on enable */
#define IRQ_REPLAY          8      /* IRQ has been replayed but not acked yet */
#define IRQ_AUTODETECT     16      /* IRQ is being autodetected */
#define IRQ_WAITING         32      /* IRQ not yet seen - for autodetection */
#define IRQ_LEVEL           64      /* IRQ level triggered */
#define IRQ_MASKED         128     /* IRQ masked - shouldn't be seen again */
#define IRQ_PER_CPU        256     /* IRQ is per CPU */

```

Registering Interrupt Service Routine

- ✓ Drivers can register an IH and enable a given interrupt line via

```
int int request_irq(unsigned int irq,
                    irqreturn_t handler,
                    unsigned long irqflags,
                    const char * devname,
                    void *dev_id);
```

- **irq**: the interrupt line # to allocate
- **handler**: pointer to actual ISR
- **irqflags**: Attributes
- **devname**: an ASCII text representation such as “keyboard”
- **dev_id**: is used as an unique cookie when this line is shared
 - A common practice is to pass driver’s device structure

irqflags Options

✓ IRQF_DISABLED

- ✓ The given IH is a fast IH: it runs with all interrupts disabled on local processor
- ✓ By default (w/o this flag), all interrupts are enabled **except** the interrupt lines of any running handlers

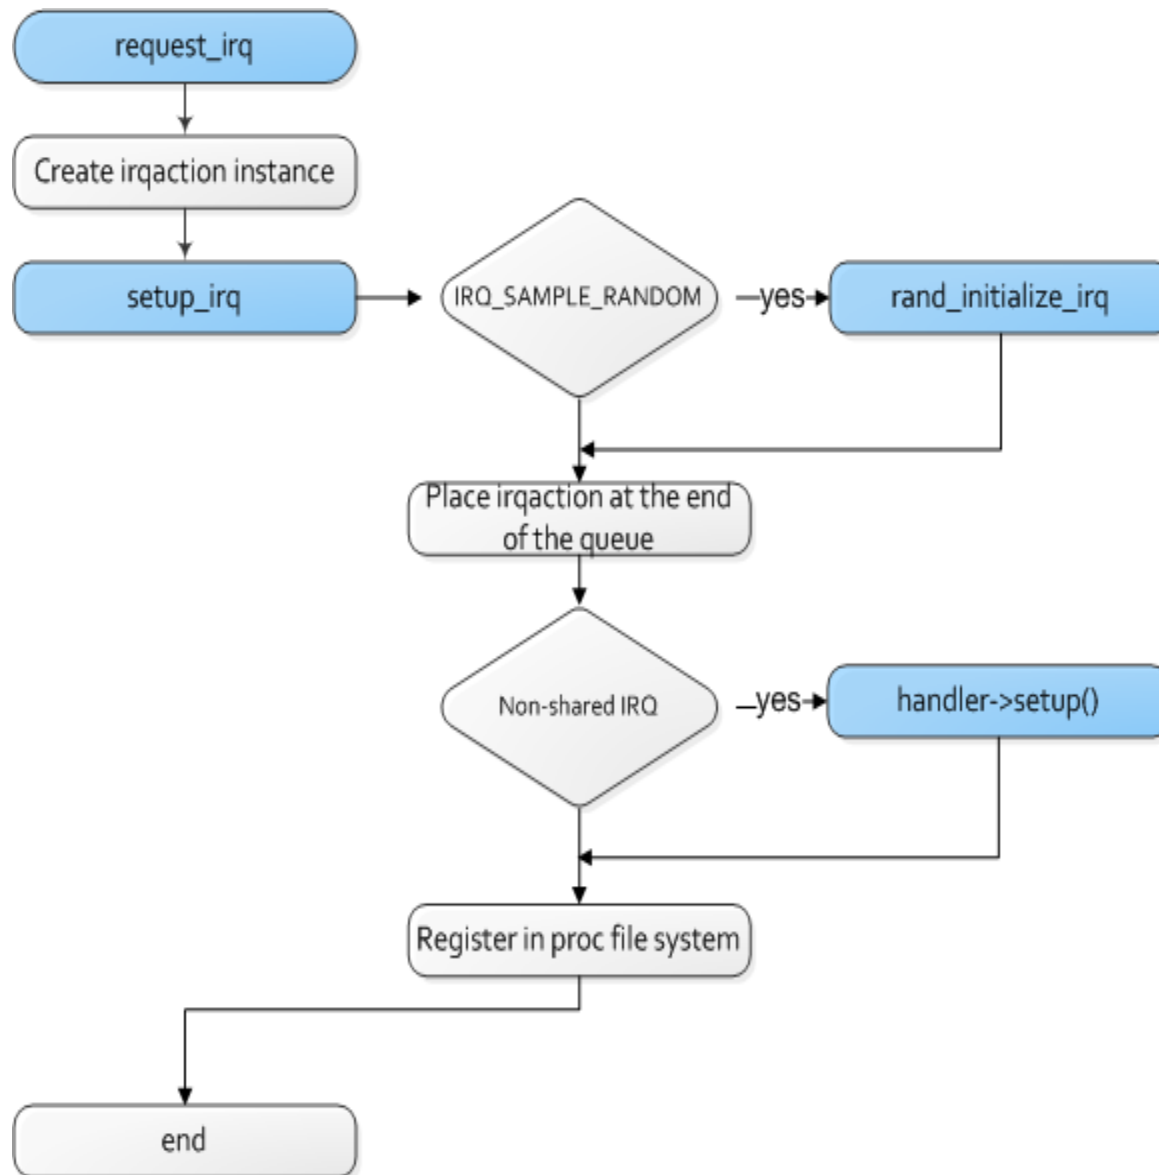
✓ IRQF_PROBE_SHARED

- ✓ Verify current irq is set to shared

✓ IRQF_SHARED

- ✓ The interrupt line can be shared among multiple ISRs
- See `include/linux/interrupt.h`

request_irq()



request_irq Usage

- ✓ To request an interrupt line and install a handler

```
if (request_irq(irqn, my_interrupt, IRQF_SHARED, "my-device", dev)) {
    printk(KERN_ERR "my_device: cannot register IRQ %d\n", irqn);
    return -EIO;
}
```

- If return 0 → handler was successfully installed

- ✓ To free an interrupt line, call

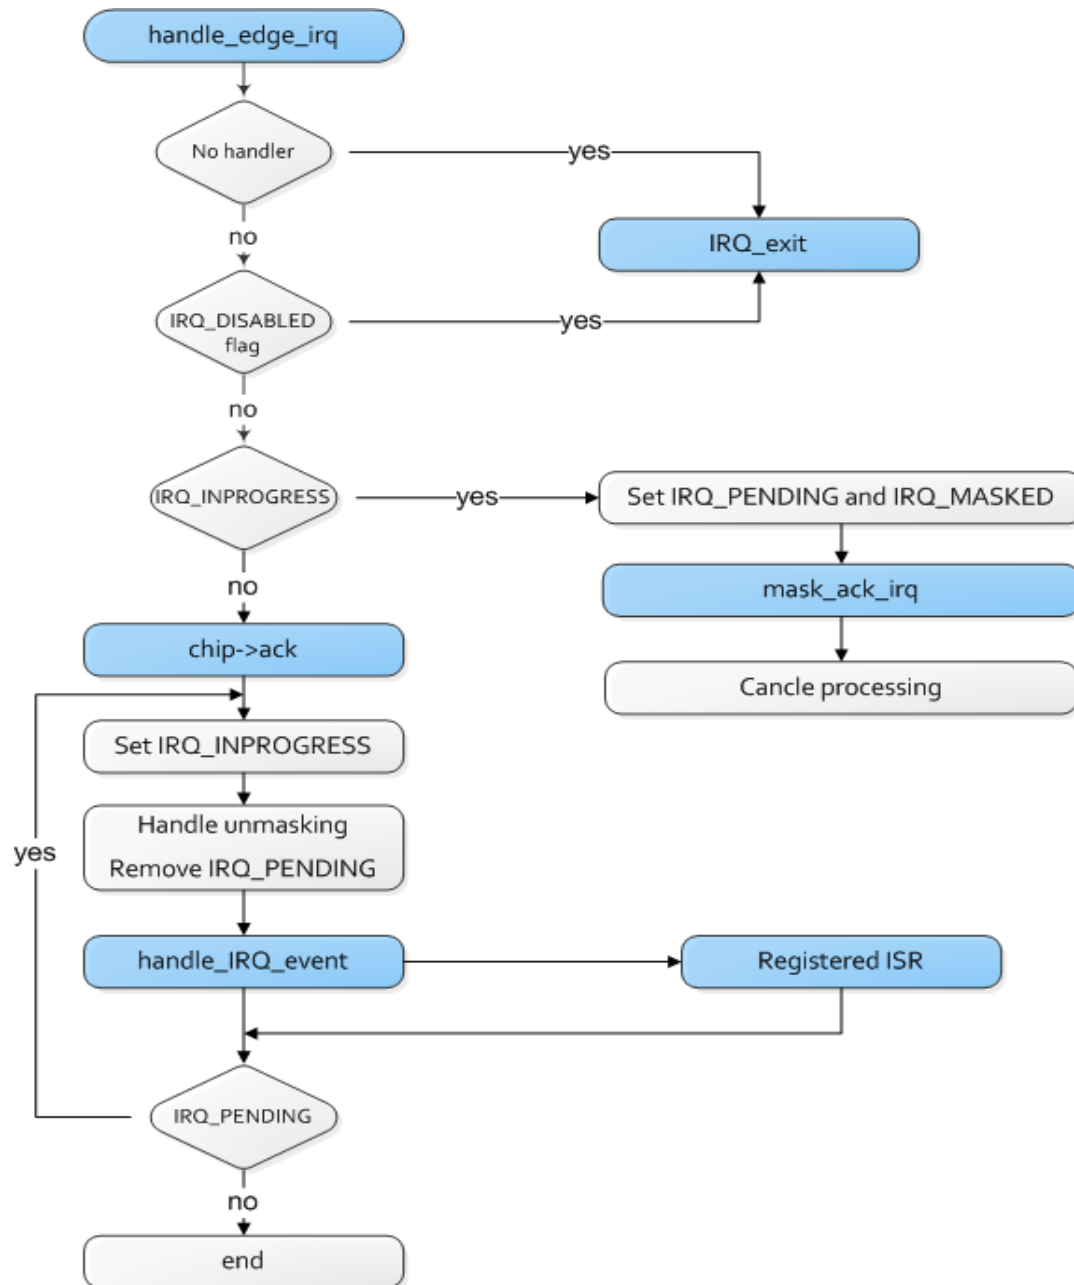
```
void free_irq(unsigned int irq, void *dev_id);
```

- If line is not shared, it removes handler and disables the line
- Otherwise, the line is only disabled at removal of last handler
- `dev_id` is used to **uniquely** identify an interrupt handler

Interrupt Flow types

- ✓ Edge-triggered interrupts are signalled by a change in the interrupt line - from low voltage to high, from high to low, or both.
- ✓ These interrupts do not necessarily have to be masked while being processed
- ✓ kernel must track "pending" interrupts, and the interrupt handler must loop until all interrupts have been dealt with.

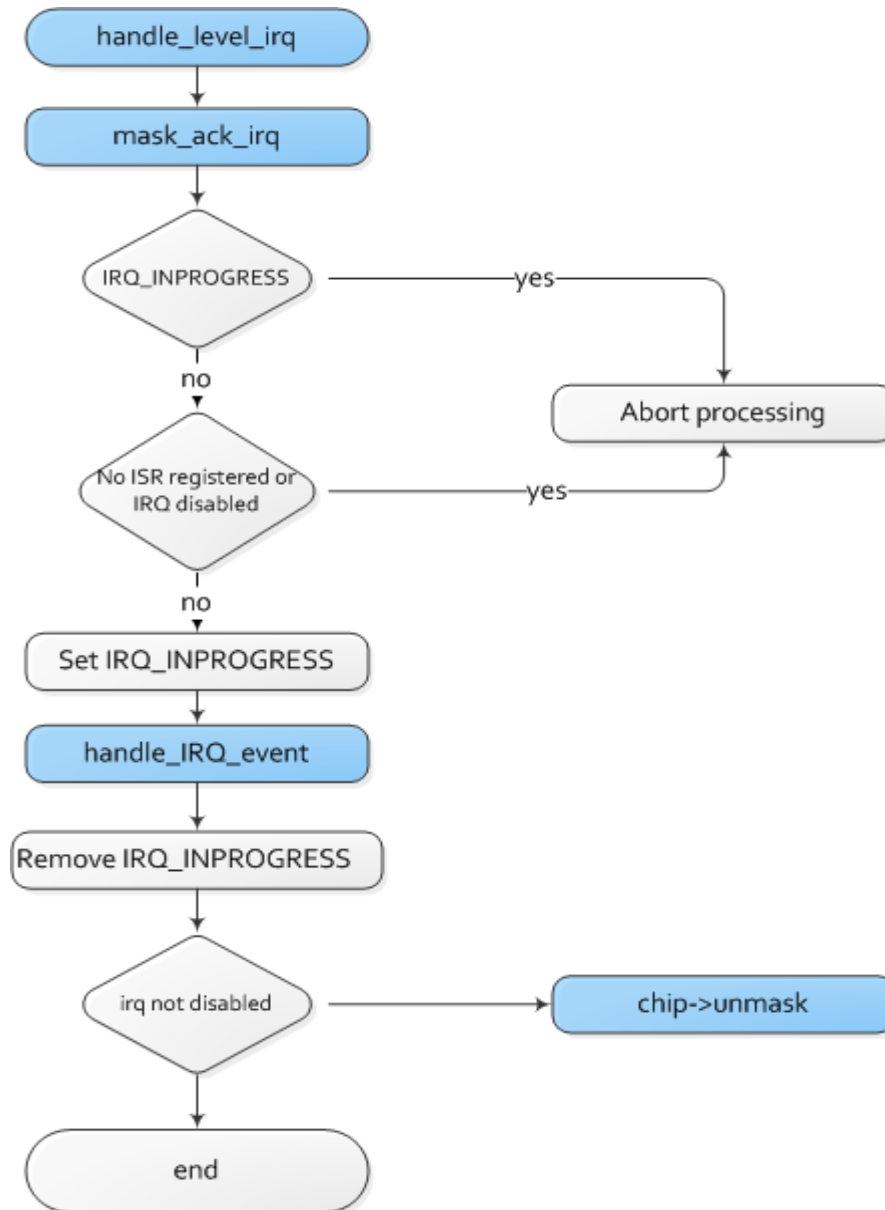
Edge handler



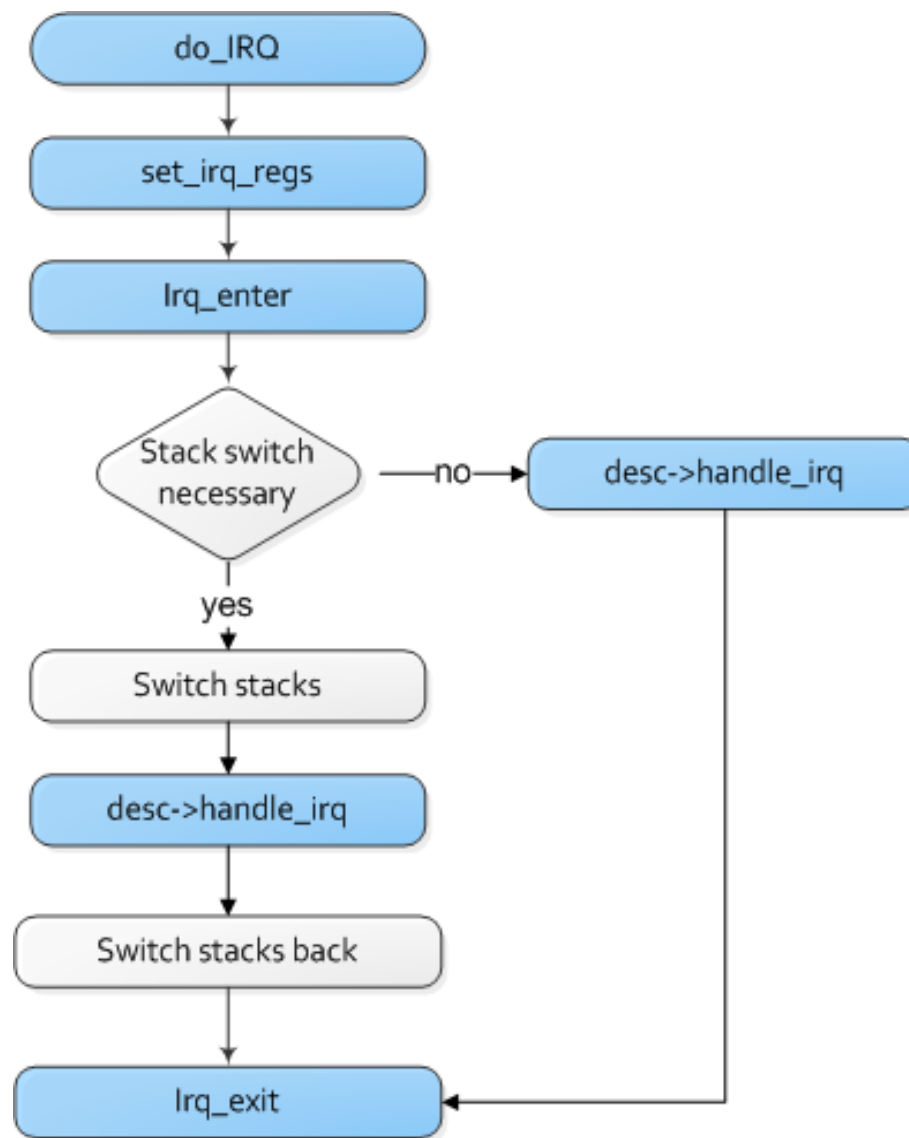
Interrupt Flow types

- ✓ A level-triggered interrupt is a class of interrupts where the presence of an interrupt is indicated by a high level (1), or low level (0), of irq line.
- ✓ A device wishing to signal an interrupt drives line to its active level, and then holds it at that level until serviced.
- ✓ **Level-triggered interrupts** are active as long as the device asserts its IRQ line.
- ✓ These interrupts must be masked while being processed

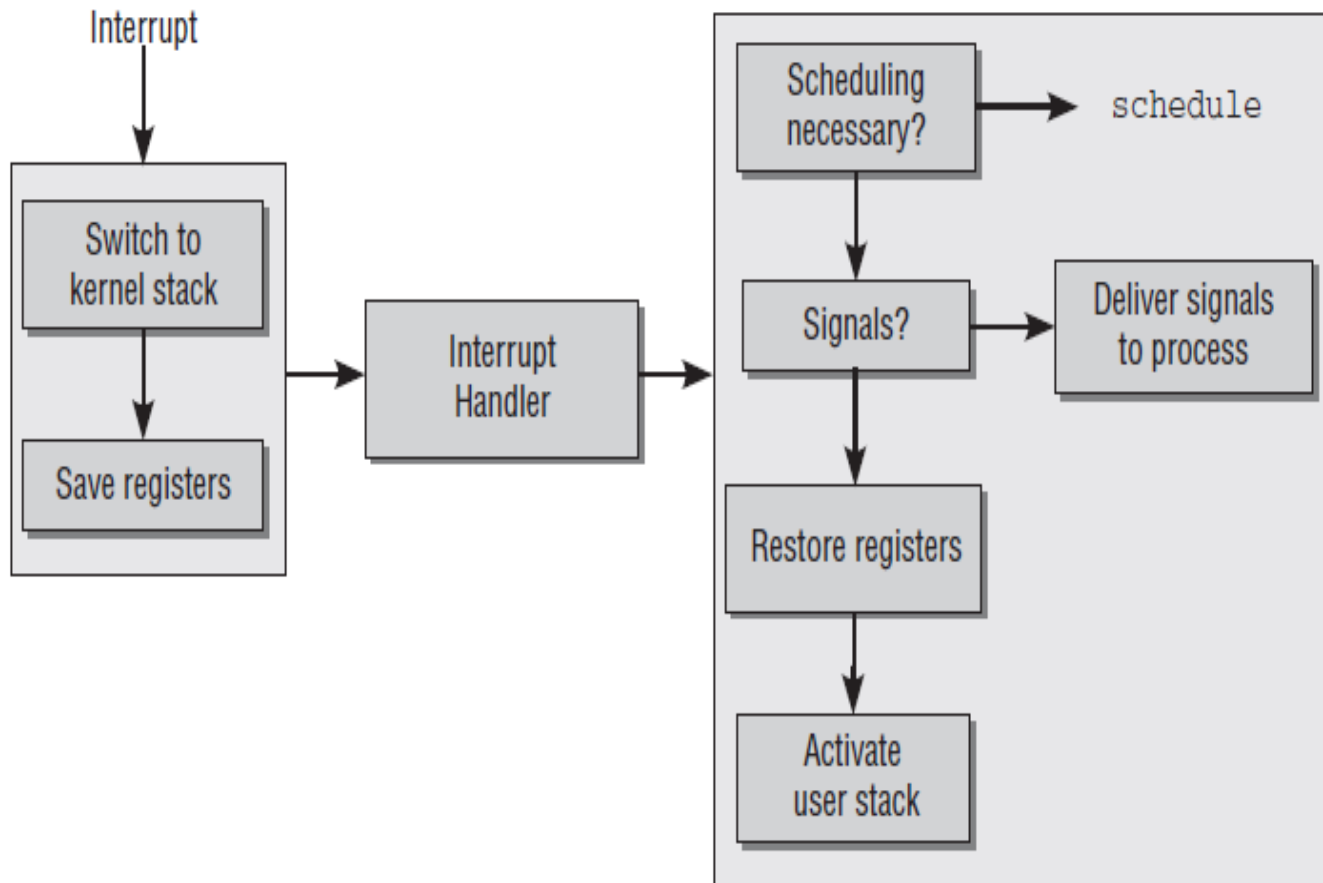
handle_level_irq



Irq stacks



Handling an Interrupt



Control Interfaces

- ✓ Purpose: to allow disabling the interrupt system for current CPU or mask out an interrupt line for entire machine
- ✓ Disable/enable interrupts locally for current processor:
 - `local_irq_disable();`
 - `local_irq_enable();`
 - `local_irq_save(flags);` // save and disable
 - `local_irq_restore(flags);` // restore and enable

Control Interfaces (2)

- ✓ Disable only a specific interrupt line for entire system
 - `disable_irq(unsigned int irq);`
 - Wait until any currently executing handler completes
 - `disable_irq_nosync(unsigned int irq);`
 - Will not wait
 - `enable_irq(unsigned int irq);`
 - If `disable_irq()` is called twice, only the 2nd `enable_irq()` will actually enable the interrupt line
 - `synchronize_irq(unsigned int irq);`
 - Wait for a specific IH to exit, if executing, before returning
- ✓ Status checking
 - `irqs_disabled()`
 - returns nonzero if interrupt system on local CPU is disabled, or 0 otherwise
 - `in_interrupt()`
 - return nonzero if kernel is in interrupt context (including in IH or BH)
 - return zero if kernel is in process context
 - `in_irq()`
 - return nonzero if kernel is executing an interrupt handler