

# Interrupt service Routines

# Introduction

- A typical declaration of an ISR
 

```
static irqreturn_t intr_handler(int irq, void *dev_id, struct pt_regs *regs)
```

  - **irq**: the IRQ line it is servicing
  - **dev\_id**: a generic pointer to the same dev\_id given to request\_irq()
  - **regs**: processor registers prior to servicing the interrupt
- Return value
  - **IRQ\_NONE**: ISR detects an interrupt for which its device was not the originator
  - **IRQ\_HANDLED**: Otherwise
- At a minimum, most ISRs need to provide acks to the device that they received the interrupt
- When a line is **shared** by multiple ISRs, kernel invokes sequentially each registered handler
  - A HW device should have a **status** register its ISR can check

# Why Bottom Half?

- IH (top halves) have following properties (requirements)
  - IH (top half) need to run as **quickly** as possible
  - IH runs with some (or all) interrupt levels **disabled**
  - IH are often time-critical and they deal with **HW**
  - IH do not run in process context and **cannot block**
- No hard and fast rules exist about what work to perform where
  - **Research work needed**
- Bottom halves are to defer work **later**
  - “Later” is often simply “**not now**”
  - Often, bottom halves run immediately after interrupt returns
  - They run with all interrupts **enabled**

# ISR Design Considerations

1. Critical actions
  - ✓ must be executed immediately following an interrupt.
  - ✓ Irq's disabled while execution.
2. Noncritical actions:
  - ✓ should be performed as quickly as possible but with enabled interrupts (they may therefore be interrupted by other system events).
3. Deferrable actions:
  - ✓ not important and need not be implemented in the interrupt handler.
  - ✓ These actions can be delayed until kernel has nothing better to do.

# Linux Bottom Halves

- Multiple mechanisms are available for implementing a bottom half
  - **softirq, tasklet, work queues**
- **softirq**: (available since 2.3)
  - A set of **32** statically defined bottom halves that can run simultaneously on any processor
    - Even 2 of the same type can run concurrently
  - Used when **performance** is critical
  - Must be registered statically at **compile-time**
- **tasklet**: (available since 2.3)
  - Are built on top of softirqs
  - Two different tasklets can run simultaneously on different processors
    - But 2 of the same type cannot run simultaneously
  - Used most of the time for its ease and flexibility
  - Code can **dynamically** register tasklets
- **work queues**: (available since 2.5)
  - Queueing work to later be performed in process context

# Softirqs

- Softirqs are rarely used
  - can concurrently run
  - Statically allocated at compile-time
- Related code: kernel/softirq.c

```
struct softirq_action
{
    void (*action)(struct softirq_action *); // function to run
    void *data; // data to pass to function
};
static struct softirq_action softirq_vec[32];
```

- In 2.6.30 kernel, only 9 softirqs are used

```
enum
{
    HI_SOFTIRQ=0,          TIMER_SOFTIRQ,
    NET_TX_SOFTIRQ,        NET_RX_SOFTIRQ,
    BLOCK_SOFTIRQ,         TASKLET_SOFTIRQ
};
```

# Using Softirqs

- Index assignment:
  - Before using softirqs, you must declare its **index** at compile time via an **enum**
  - Softirqs with **lower** numerical priority execute **first**
- Register handler:
  - Softirq handler is registered at run-time via `open_softirq()`

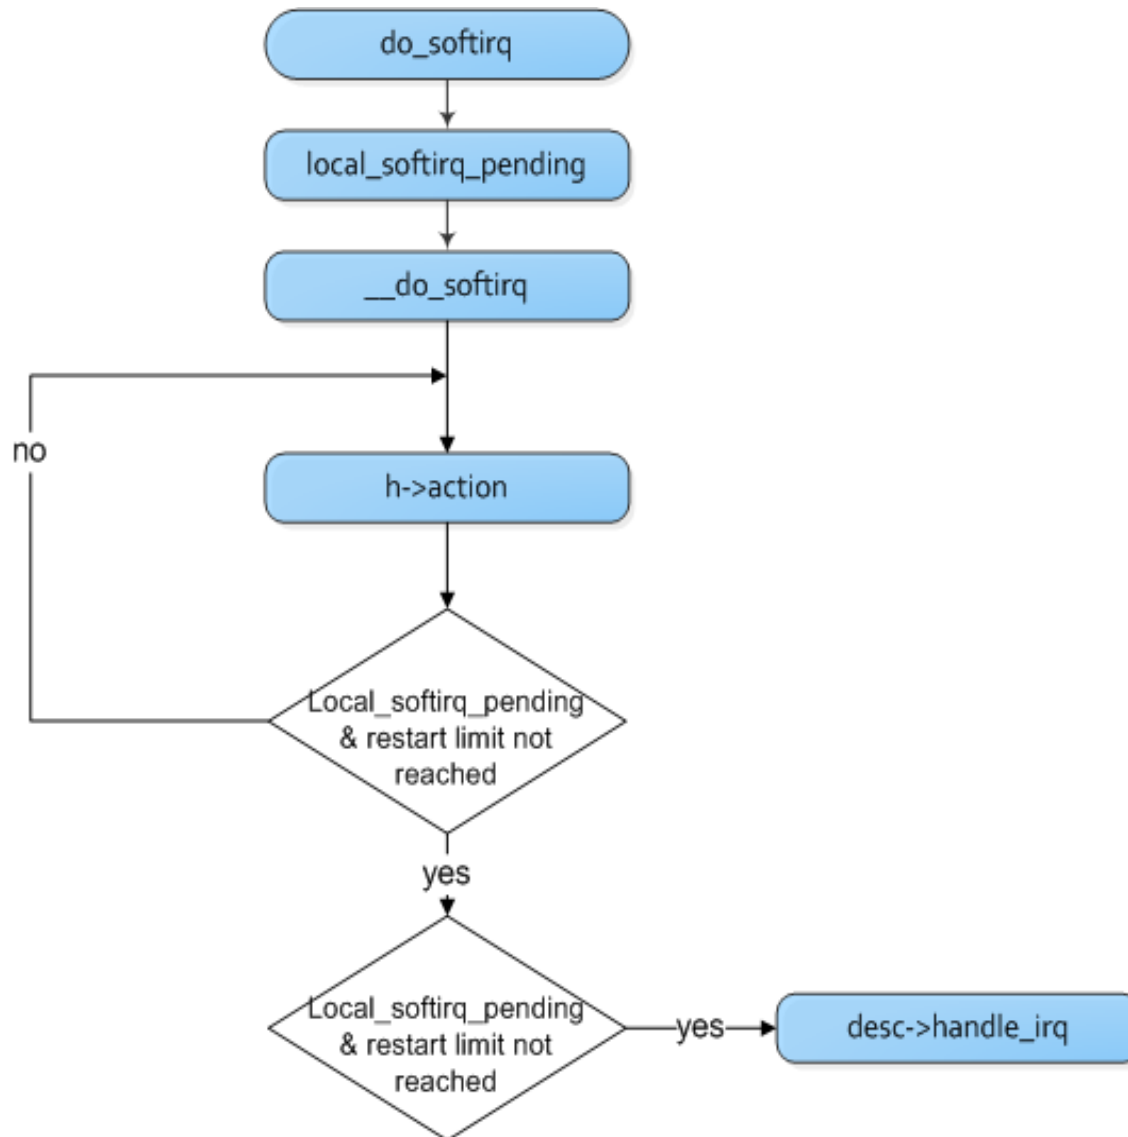
```
void open_softirq(int nr, void (*action)(struct softirq_action*), void *data)
{
    softirq_vec[nr].data = data;
    softirq_vec[nr].action = action;
}
```

# Using Softirq (2/2)

- Raising softirq
  - Call: `raise_softirq(NEX_TX_SOFTIRQ)`, for example
  - Softirqs are often raised from within interrupt handlers
  - When done processing interrupts, kernel invokes `do_softirq()`
- Softirqs run with interrupt **enabled** and **cannot** sleep



# do\_softirq



# ksoftirqd

- Most commonly, kernel processes softirqs on return from handling an interrupt
  - In **interrupt context**
- However, softirqs may be raised at very high rates
  - Sometimes, they reactivate themselves
  - It may lead to starvation of user programs
- Kernel solution
  - When softirqs grow excessively, kernel wakes up a family of **kernel threads**
  - One thread per processor, named **ksoftirqd/n**
  - `static int ksoftirqd(void * __bind_cpu) [code]`

## Tasklet Properties

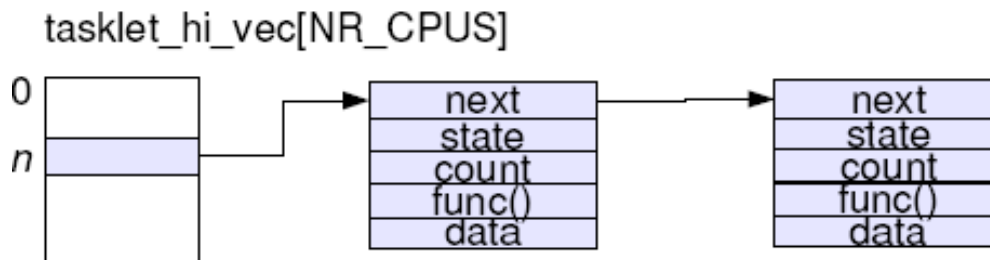
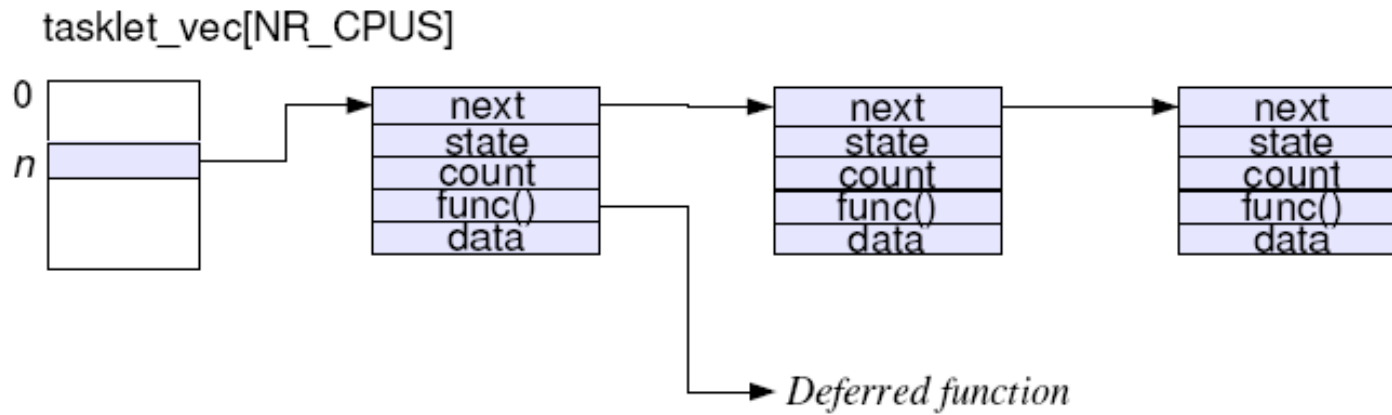
- Guaranteed to run once
  - Once scheduled, a tasklet is guaranteed to be executed once after that
  - An already scheduled but not yet executed tasklet can be rescheduled, but will be executed only once
  - Once a tasklet starts running, it can be rescheduled to run again later
  - Tasklet is strictly serialized (no nesting)
  - Different tasklets can run simultaneously on different CPUs

## Tasklet Data Structure

- Deferred function and the argument
- Data Structure: in include/linux/interrupt.h
 

```
struct tasklet_struct
{
    struct tasklet_struct *next;
    unsigned long state;
    atomic_t count;
    void (*func)(unsigned long);
    unsigned long data;
};
```
- Two tasklet lists per CPU (one higher priority)

## Tasklet Lists



## Declaring Tasklet

- Define a function that takes one argument
  - void function(unsigned long data)
- DECLARE\_TASKLET(name,function,data)
  - Declares a tasklet (of type struct tasklet\_struct) with the given name, function, and (unsigned long) data value (as the argument when function is later called).
- DECLARE\_TASKLET\_DISABLED(name,function,data)
  - Declares a tasklet but with initial state “disabled” -- it can be scheduled but will not be executed until enabled at some future time

## Using Tasklet

- To schedule a tasklet to run soon:
  - `void tasklet_schedule(struct tasklet_struct *t)`
  - `void tasklet_hi_schedule(struct tasklet_struct *t)`
    - Both functions: add to the beginning of corresponding tasklet list, and raise softirq
- Disable or enable a tasklet:
  - `void tasklet_disable(struct tasklet_struct *t);`
  - `void tasklet_enable(struct tasklet_struct *t);`
    - Disabled tasklet can be scheduled but won't run -- will remain in the tasklet list until enabled again