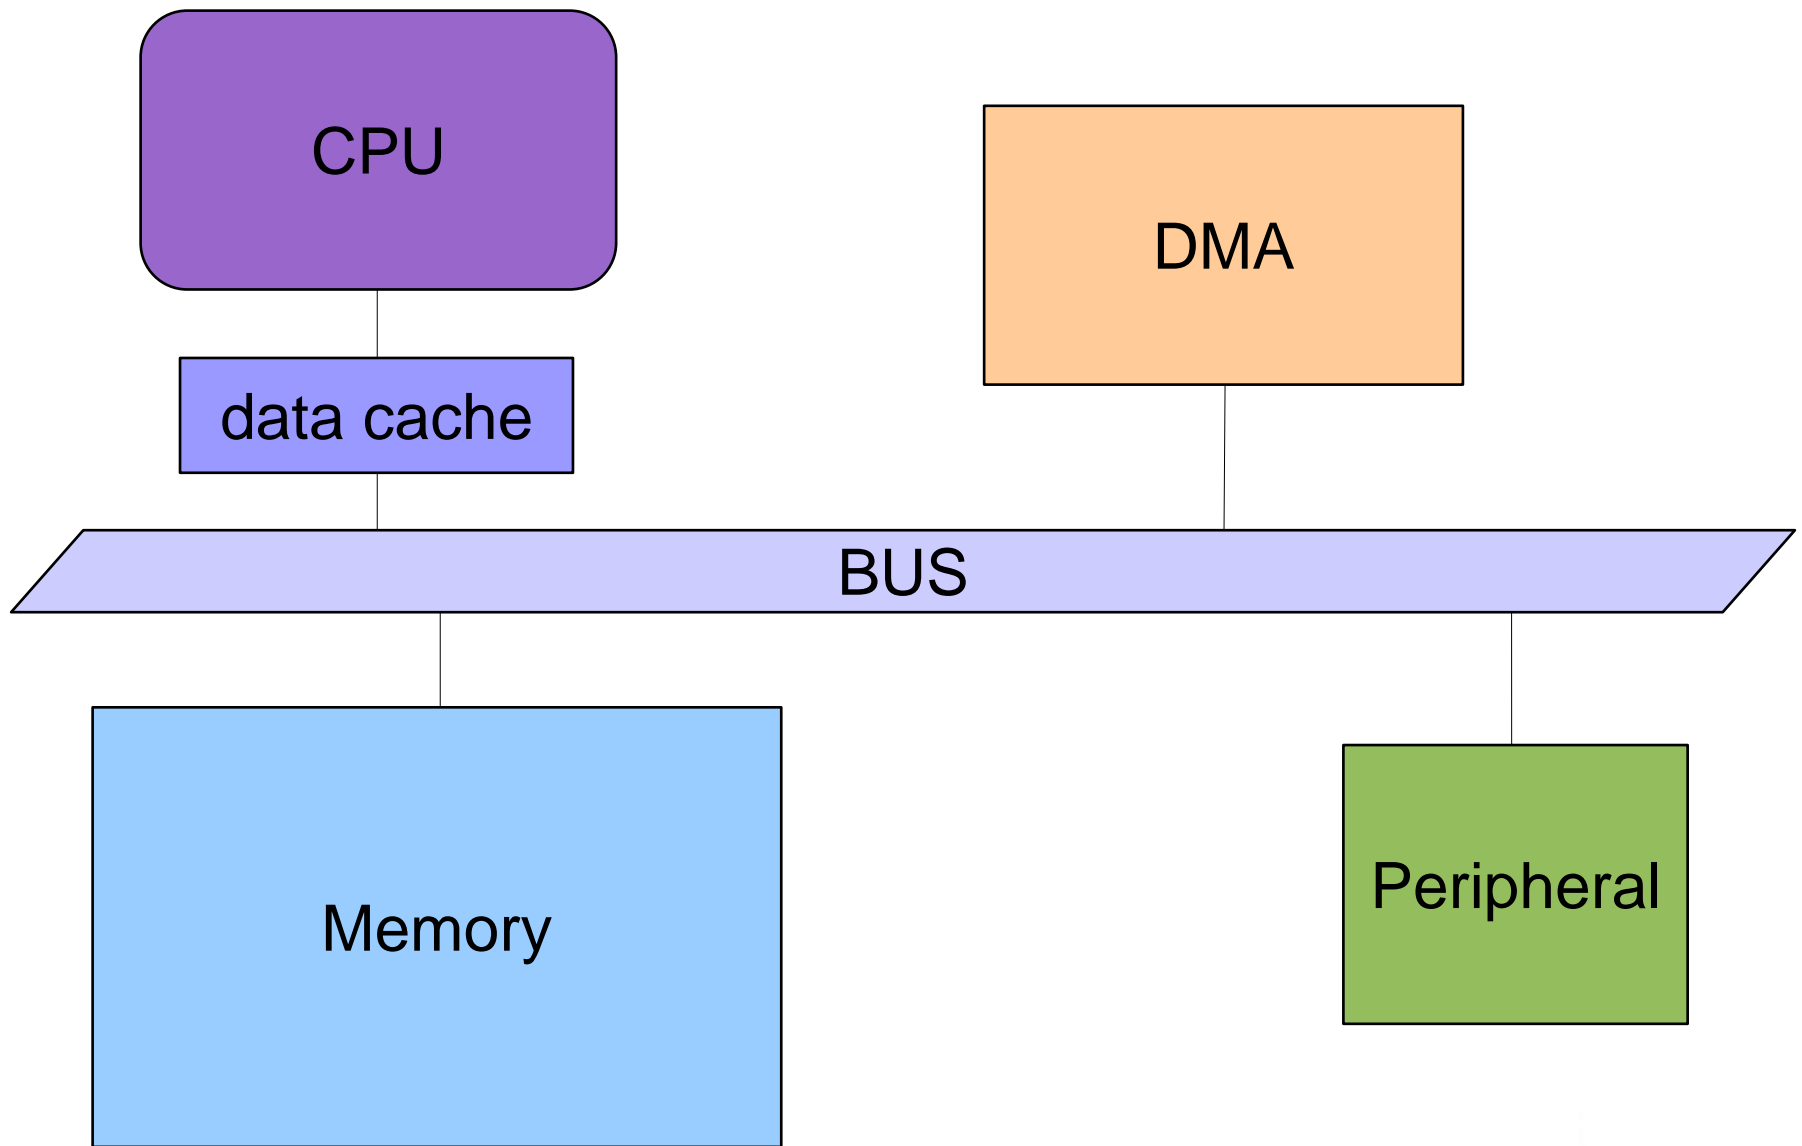


# Linux DMA

# DMA integration



# Constraints with a DMA

- ▶ A DMA deals with physical addresses, so:
  - ▶ Programming a DMA requires retrieving a physical address at some point (virtual addresses are usually used)
  - ▶ The memory accessed by the DMA shall be physically contiguous
- ▶ The CPU can access memory through a data cache
  - ▶ Using the cache can be more efficient (faster accesses to the cache than the bus)
  - ▶ But the DMA does not access to the CPU cache, so one need to take care of cache coherency (cache content vs memory content)
  - ▶ Either flush or invalidate the cache lines corresponding to the buffer accessed by DMA and processor at strategic times

# DMA memory constraints

- ▶ Need to use contiguous memory in physical space.
- ▶ Can use any memory allocated by `kmalloc` (up to 128 KB) or `__get_free_pages` (up to 8MB).
- ▶ Can use block I/O and networking buffers, designed to support DMA.
- ▶ Can not use `vmalloc` memory (would have to setup DMA on each individual physical page).

# Reserving memory for DMA

To make sure you've got enough RAM for big DMA transfers...  
Example assuming you have 32 MB of RAM, and need 2 MB for DMA:

- ▶ Boot your kernel with `mem=30`

The kernel will just use the first 30 MB of RAM.

- ▶ Driver code can now reclaim the 2 MB left:

```
dmabuf = ioremap (
    0x1e00000,          /* Start: 30 MB */
    0x200000);         /* Size: 2 MB */
```

# Memory synchronization issues

Memory caching could interfere with DMA

- ▶ Before DMA to device:  
Need to make sure that all writes to DMA buffer are committed.
- ▶ After DMA from device:  
Before drivers read from DMA buffer, need to make sure that memory caches are flushed.
- ▶ Bidirectional DMA  
Need to flush caches before and after the DMA transfer.

# Linux DMA API

The kernel DMA utilities can take care of:

- ▶ Either allocating a buffer in a cache coherent area,
- ▶ Or making sure caches are flushed when required,
- ▶ Managing the DMA mappings and IOMMU (if any).
- ▶ See Documentation/DMA-API.txt  
for details about the Linux DMA generic API.
- ▶ Most subsystems (such as PCI or USB) supply their own DMA API, derived from the generic one. May be sufficient for most needs.

# Coherent or streaming DMA mappings

## Coherent mappings

The kernel allocates a suitable buffer and sets the mapping for the driver.

- ▶ Can simultaneously be accessed by the CPU and device.
- ▶ So, has to be in a cache coherent memory area.
- ▶ Usually allocated for the whole time the module is loaded.
- ▶ Can be expensive to setup and use on some platforms.

## Streaming mappings

The kernel just sets the mapping for a buffer provided by the driver.

- ▶ Use a buffer already allocated by the driver.
- ▶ Mapping set up for each transfer. Keeps DMA registers free on the hardware.
- ▶ Some optimizations also available.
- ▶ The recommended solution.





# Allocating coherent mappings

The kernel takes care of both buffer allocation and mapping:

```
include <asm/dma-mapping.h>
```

```
void *                                /* Output: buffer address */
dma_alloc_coherent(
    struct device *dev,               /* device structure */
    size_t size,                     /* Needed buffer size in bytes */
    dma_addr_t *handle,              /* Output: DMA bus address */
    gfp_t gfp                        /* Standard GFP flags */
);
```

```
void dma_free_coherent(struct device *dev,
    size_t size, void *cpu_addr, dma_addr_t handle);
```



# Setting up streaming mappings

Works on buffers **already allocated by the driver**

```
<include linux/dmapool.h>
```

```
dma_addr_t dma_map_single(  
    struct device *,                                /* device structure */  
    void *,                                          /* input: buffer to use */  
    size_t,                                          /* buffer size */  
    enum dma_data_direction /* Either DMA_BIDIRECTIONAL,  
                                DMA_TO_DEVICE or  
                                DMA_FROM_DEVICE */  
);
```

```
void dma_unmap_single(struct device *dev, dma_addr_t handle, size_t  
    size, enum dma_data_direction dir);
```

# DMA streaming mapping notes

- ▶ When the mapping is active: only the device should access the buffer (potential cache issues otherwise).
- ▶ The CPU can access the buffer only after unmapping! Use locking to prevent CPU access to the buffer.
- ▶ Another reason: if required, this API can create an intermediate *bounce buffer* (used if the given buffer is not usable for DMA).
- ▶ The Linux API also supports scatter / gather DMA streaming mappings.

# DMA summary

Most drivers can use the specific API provided by their subsystem: USB, PCI, SCSI... Otherwise they can use the Linux generic API:

## Coherent mappings

- ▶ DMA buffer allocated by the kernel
- ▶ Set up for the whole module life
- ▶ Can be expensive. Not recommended.
- ▶ Let both the CPU and device access the buffer at the same time.
- ▶ Main functions:  
`dma_alloc_coherent`  
`dma_free_coherent`

## Streaming mappings

- ▶ DMA buffer allocated by the driver
- ▶ Set up for each transfer
- ▶ Cheaper. Saves DMA registers.
- ▶ Only the device can access the buffer when the mapping is active.
- ▶ Main functions:  
`dma_map_single`  
`dma_unmap_single`